

Business Problem

The main aim of analyzing the backorder data is to find out the reason due to which a particular product is not in stock when a potential customer wished to buy it. Therefore, the companies are looking for to explore the cause of backorder and the solution which might be used to minimize the backorder and increase the business. There are several causes that could be the reason of backorder

- Delayed in placing the order - Based on the ordercycle, stock etc. orders are placed to the vendor daily, weekly or even it takes a month or some time interval. Even after placing the order an individual person reviewed the order and came out with a conclusion that whether the particular order is need to be executed or not. For a delay in decision making results a backorder.
- Warehouse Discrepancies - A discrepancies may occure when the stock maintained digitally or by some mannual measure is not matched with the actual stock quantity present in warehouse.
- On other hand human error which is caused by some issue of an individual person can cause a backorder.
- If the production in factories is lacking due to their internal issue, which no E-commerce can control by their own can lead to backorder.
- High Demand caused due to large numbers of order placement by customers. This can be a result for customer purchasing in an abnormal manner or for some seasonal demand like festive season. For this type of situation, we must find out the reason behind this demand.

ML Formulation

Identifying products with the highest chances of shortage prior its occurrence can present a high opportunity to improve an overall company's performance. Machine learning is applied on the design and development of predictive models which assess all areas of management, providing essential insights for companies to understand and take action to changes in its operation.

Accuracy metric

- **ROC-AUC:** The Receiver Operator Characteristic (ROC) curve is an evaluation metric for binary classification problems. It is a probability curve that plots the TPR against FPR at various threshold values. The Area Under the Curve (AUC) is the measure of the ability of a classifier to distinguish between classes and is used as a summary of the ROC curve. The higher the AUC, the better the performance of the model at distinguishing between the positive and negative classes. So, if for a model say $m1$ $AUC > m2$ AUC then at most of the threshold values model $m1$ is able to identify the positive class better than negative class.
- **Precision Recall AUC:** Precision is measured as out of the total number of predicted positive points percentage of actual positive prediction.
Mathematically $Pr = TP / (TP + FP)$.
On the other hand recall is measured as the percentage of actual positive prediction out of the total number of positive points in the dataset.
Mathematically $Rc = TP / (TP + FN)$.
Both the precision and the recall are focused on the positive class (the minority class) and are unconcerned with the true negatives (majority class). As it is an imbalance dataset the precision and recall make it possible to assess the performance of a classifier on the minority class. AUC of Pr-Re curve is very important as the business needs to select the suitable threshold based on treade off b/w precision and recall.

Data Source

The data is taken from the github of the following link

https://github.com/rodrigasantis1/backorder_prediction which stated that the main source of data is from Kaggle's "Can You Predict Product Backorders?"

In [1]:

```
!pip3 install tensorflow
```

```
Requirement already satisfied: tensorflow in /opt/conda/lib/python3.7/site-packages (2.4.1)
Requirement already satisfied: grpcio~=1.32.0 in /opt/conda/lib/python3.7/site-packages (from tensorflow) (1.32.0)
Requirement already satisfied: wheel~=0.35 in /opt/conda/lib/python3.7/site-packages (from tensorflow) (0.36.2)
Requirement already satisfied: tensorboard~=2.4 in /opt/conda/lib/python3.7/site-packages (from tensorflow) (2.5.0)
Requirement already satisfied: h5py~=2.10.0 in /opt/conda/lib/python3.7/site-packages (from tensorflow) (2.10.0)
Requirement already satisfied: six~=1.15.0 in /opt/conda/lib/python3.7/site-packages (from tensorflow) (1.15.0)
Requirement already satisfied: wrapt~=1.12.1 in /opt/conda/lib/python3.7/site-packages (from tensorflow) (1.12.1)
Requirement already satisfied: typing-extensions~=3.7.4 in /opt/conda/lib/python3.7/site-packages (from tensorflow) (3.7.4.3)
Requirement already satisfied: termcolor~=1.1.0 in /opt/conda/lib/python3.7/site-packages (from tensorflow) (1.1.0)
Requirement already satisfied: astunparse~=1.6.3 in /opt/conda/lib/python3.7/site-packages (from tensorflow) (1.6.3)
Requirement already satisfied: tensorflow-estimator<2.5.0,>=2.4.0 in /opt/conda/lib/python3.7/site-packages (from tensorflow) (2.4.0)
Requirement already satisfied: gast==0.3.3 in /opt/conda/lib/python3.7/site-packages (from tensorflow) (0.3.3)
Requirement already satisfied: absl-py~=0.10 in /opt/conda/lib/python3.7/site-packages (from tensorflow) (0.12.0)
Requirement already satisfied: keras-preprocessing~=1.1.2 in /opt/conda/lib/python3.7/site-packages (from tensorflow) (1.1.2)
Requirement already satisfied: google-pasta~=0.2 in /opt/conda/lib/python3.7/site-packages (from tensorflow) (0.2.0)
Requirement already satisfied: protobuf>=3.9.2 in /opt/conda/lib/python3.7/site-packages (from tensorflow) (3.15.8)
Requirement already satisfied: numpy~=1.19.2 in /opt/conda/lib/python3.7/site-packages (from tensorflow) (1.19.5)
Requirement already satisfied: opt-einsum~=3.3.0 in /opt/conda/lib/python3.7/site-packages (from tensorflow) (3.3.0)
Requirement already satisfied: flatbuffers~=1.12.0 in /opt/conda/lib/python3.7/site-packages (from tensorflow) (1.12)
Requirement already satisfied: google-auth<2,>=1.6.3 in /opt/conda/lib/python3.7/site-packages (from tensorflow) (1.28.0)
Requirement already satisfied: google-auth-oauthlib<0.5,>=0.4.1 in /opt/conda/lib/python3.7/site-packages (from tensorflow) (0.4.3)
Requirement already satisfied: werkzeug>=0.11.15 in /opt/conda/lib/python3.7/site-packages (from tensorflow) (1.0.1)
Requirement already satisfied: tensorboard-data-server<0.7.0,>=0.6.0 in /opt/conda/lib/python3.7/site-packages (from tensorflow) (0.6.0)
Requirement already satisfied: requests<3,>=2.21.0 in /opt/conda/lib/python3.7/site-packages (from tensorflow) (2.25.1)
Requirement already satisfied: setuptools>=41.0.0 in /opt/conda/lib/python3.7/site-packages (from tensorflow) (49.6.0.post20210108)
Requirement already satisfied: markdown>=2.6.8 in /opt/conda/lib/python3.7/site-packages (from tensorflow) (3.3.4)
Requirement already satisfied: tensorboard-plugin-wit>=1.6.0 in /opt/conda/lib/python3.7/site-packages (from tensorflow) (1.8.0)
Requirement already satisfied: pyasn1-modules>=0.2.1 in /opt/conda/lib/python3.7/site-packages (from tensorflow) (0.2.7)
Requirement already satisfied: cachetools<5.0,>=2.0.0 in /opt/conda/lib/python3.7/site-packages (from tensorflow) (4.2.1)
Requirement already satisfied: rsa<5,>=3.1.4 in /opt/conda/lib/python3.7/site-packages (from tensorflow) (4.7.2)
Requirement already satisfied: requests-oauthlib>=0.7.0 in /opt/conda/lib/python3.7/site-packages (from tensorflow) (1.3.0)
Requirement already satisfied: importlib-metadata in /opt/conda/lib/python3.7/site-packages (from tensorflow) (3.10.1)
```

Requirement already satisfied: pyasn1<0.5.0,>=0.4.6 in /opt/conda/lib/python3.7/site-packages (from pyasn1-modules>=0.2.1->google-auth<2,>=1.6.3->tensorboard~=2.4->tensorflow) (0.4.8)
Requirement already satisfied: urllib3<1.27,>=1.21.1 in /opt/conda/lib/python3.7/site-packages (from requests<3,>=2.21.0->tensorboard~=2.4->tensorflow) (1.26.4)
Requirement already satisfied: idna<3,>=2.5 in /opt/conda/lib/python3.7/site-packages (from requests<3,>=2.21.0->tensorboard~=2.4->tensorflow) (2.10)
Requirement already satisfied: chardet<5,>=3.0.2 in /opt/conda/lib/python3.7/site-packages (from requests<3,>=2.21.0->tensorboard~=2.4->tensorflow) (4.0.0)
Requirement already satisfied: certifi>=2017.4.17 in /opt/conda/lib/python3.7/site-packages (from requests<3,>=2.21.0->tensorboard~=2.4->tensorflow) (2020.12.5)
Requirement already satisfied: oauthlib>=3.0.0 in /opt/conda/lib/python3.7/site-packages (from requests-oauthlib>=0.7.0->google-auth-oauthlib<0.5,>=0.4.1->tensorboard~=2.4->tensorflow) (3.0.1)
Requirement already satisfied: zipp>=0.5 in /opt/conda/lib/python3.7/site-packages (from importlib-metadata->markdown>=2.6.8->tensorboard~=2.4->tensorflow) (3.4.1)

In [2]: `!pip3 install keras`

Requirement already satisfied: keras in /opt/conda/lib/python3.7/site-packages (2.4.3)
Requirement already satisfied: h5py in /opt/conda/lib/python3.7/site-packages (from keras) (2.10.0)
Requirement already satisfied: scipy>=0.14 in /opt/conda/lib/python3.7/site-packages (from keras) (1.6.2)
Requirement already satisfied: pyyaml in /opt/conda/lib/python3.7/site-packages (from keras) (5.4.1)
Requirement already satisfied: numpy>=1.9.1 in /opt/conda/lib/python3.7/site-packages (from keras) (1.19.5)
Requirement already satisfied: six in /opt/conda/lib/python3.7/site-packages (from h5py->keras) (1.15.0)

In [3]: `!pip3 install imblearn`

Requirement already satisfied: imblearn in /opt/conda/lib/python3.7/site-packages (0.0)
Requirement already satisfied: imbalanced-learn in /opt/conda/lib/python3.7/site-packages (from imblearn) (0.8.0)
Requirement already satisfied: numpy>=1.13.3 in /opt/conda/lib/python3.7/site-packages (from imbalanced-learn->imblearn) (1.19.5)
Requirement already satisfied: joblib>=0.11 in /opt/conda/lib/python3.7/site-packages (from imbalanced-learn->imblearn) (1.0.1)
Requirement already satisfied: scikit-learn>=0.24 in /opt/conda/lib/python3.7/site-packages (from imbalanced-learn->imblearn) (0.24.1)
Requirement already satisfied: scipy>=0.19.1 in /opt/conda/lib/python3.7/site-packages (from imbalanced-learn->imblearn) (1.6.2)
Requirement already satisfied: threadpoolctl>=2.0.0 in /opt/conda/lib/python3.7/site-packages (from scikit-learn>=0.24->imbalanced-learn->imblearn) (2.1.0)

In [4]: `!pip3 install xgboost`

Requirement already satisfied: xgboost in /opt/conda/lib/python3.7/site-packages (1.4.1)
Requirement already satisfied: numpy in /opt/conda/lib/python3.7/site-packages (from xgboost) (1.19.5)
Requirement already satisfied: scipy in /opt/conda/lib/python3.7/site-packages (from xgboost) (1.6.2)

In [5]: `import warnings
warnings.filterwarnings('ignore')`

```
import pandas as pd  
import numpy as np  
import seaborn as sns  
import scipy.stats as stats
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.experimental import enable_iterative_imputer  
from sklearn.impute import IterativeImputer  
from sklearn.preprocessing import RobustScaler  
from sklearn.preprocessing import StandardScaler
```

```
# from sklearn.model_selection import GridSearchCV
```

```

from sklearn.model_selection import RandomizedSearchCV
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from imblearn.ensemble import BalancedBaggingClassifier
from xgboost import XGBClassifier
from sklearn.ensemble import AdaBoostClassifier
from sklearn.utils import shuffle

from sklearn.metrics import precision_score
from sklearn.metrics import recall_score
from sklearn.metrics import roc_auc_score
from sklearn.metrics import f1_score
from sklearn.metrics import auc
from sklearn.metrics import confusion_matrix

%matplotlib inline
from matplotlib import pyplot as plot

```


```
In [6]: train_df = pd.read_csv('train.csv')
```

```
In [7]: train_df.head()
```

```
Out[7]:
```

	sku	national_inv	lead_time	in_transit_qty	forecast_3_month	forecast_6_month	forecast_9_month	sales_
0	1026827	0.0	NaN	0.0	0.0	0.0	0.0	
1	1043384	2.0	9.0	0.0	0.0	0.0	0.0	
2	1043696	2.0	NaN	0.0	0.0	0.0	0.0	
3	1043852	7.0	8.0	0.0	0.0	0.0	0.0	
4	1044048	8.0	NaN	0.0	0.0	0.0	0.0	

5 rows × 23 columns



Shape of training data

```
In [8]: train_df.shape
```

```
Out[8]: (1687861, 23)
```

The dataset contain **1687861 number of rows (records)** and **23 number of cloumns (features)**.

The dataset contain total 23 number of columns which are as follows

```
In [9]: train_df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1687861 entries, 0 to 1687860
Data columns (total 23 columns):
#   Column                Non-Null Count  Dtype
---  -
0   sku                   1687861 non-null object
1   national_inv          1687860 non-null float64
2   lead_time             1586967 non-null float64
3   in_transit_qty        1687860 non-null float64
4   forecast_3_month      1687860 non-null float64
5   forecast_6_month      1687860 non-null float64
6   forecast_9_month      1687860 non-null float64
7   sales_1_month         1687860 non-null float64

```

```

8  sales_3_month      1687860 non-null float64
9  sales_6_month      1687860 non-null float64
10 sales_9_month      1687860 non-null float64
11 min_bank           1687860 non-null float64
12 potential_issue    1687860 non-null object
13 pieces_past_due    1687860 non-null float64
14 perf_6_month_avg   1687860 non-null float64
15 perf_12_month_avg  1687860 non-null float64
16 local_bo_qty       1687860 non-null float64
17 deck_risk          1687860 non-null object
18 oe_constraint       1687860 non-null object
19 ppap_risk          1687860 non-null object
20 stop_auto_buy      1687860 non-null object
21 rev_stop           1687860 non-null object
22 went_on_backorder  1687860 non-null object
dtypes: float64(15), object(8)
memory usage: 296.2+ MB

```

The columns are of **2 different** datatypes viz - **1. Object** and **2. float64**. So we can say that the columns which have datatype of float64 are contain **numerical features (total 15 no. of numerical features)** and the columns which have datatype of object are **contain categorical features (total 8 no. of categorical features)**.

The dataset columns contain following data

- **sku** – Stock Keeping Unit for the product (In is actually a unique id for every variation of the product)
- **national_inv** - Current inventory level for the part
- **lead_time** - Transit time for product (if available)
- **in_transit_qty** - Amount of product in transit from source
- **forecast_3_month** - Forecast sales for the next 3 months
- **forecast_6_month** - Forecast sales for the next 6 months
- **forecast_9_month** - Forecast sales for the next 9 months
- **sales_1_month** - Sales quantity for the prior 1 month time period
- **sales_3_month** - Sales quantity for the prior 3 month time period
- **sales_6_month** - Sales quantity for the prior 6 month time period
- **sales_9_month** - Sales quantity for the prior 9 month time period
- **min_bank** - Minimum recommended amount to stock
- **potential_issue** - Source issue for part identified
- **pieces_past_due** - Parts overdue from source
- **perf_6_month_avg** - Source performance for prior 6 month period
- **perf_12_month_avg** - Source performance for prior 12 month period
- **local_bo_qty** - Amount of stock orders overdue
- **deck_risk** - Part risk flag
- **oe_constraint** - Part risk flag
- **ppap_risk** - Part risk flag
- **stop_auto_buy** - Part risk flag
- **rev_stop** - Part risk flag
- **went_on_backorder** - Product actually went on backorder. This is the target value.

Separating the list of categorical and numerical features

In [10]:

```

categorical_features = []
numerical_features = []
for col in train_df.columns:
    if (train_df.dtypes[col] == 'object'):
        categorical_features.append(col)
    else:
        numerical_features.append(col)
print(categorical_features)
print(numerical_features)

```

```
['sku', 'potential_issue', 'deck_risk', 'oe_constraint', 'ppap_risk', 'stop_auto_buy', 'rev_stop', 'went_on_backorder']
['national_inv', 'lead_time', 'in_transit_qty', 'forecast_3_month', 'forecast_6_month', 'forecast_9_month', 'sales_1_month', 'sales_3_month', 'sales_6_month', 'sales_9_month', 'min_bank', 'pieces_past_due', 'perf_6_month_avg', 'perf_12_month_avg', 'local_bo_qty']
```

For better data visualisation we had split the **columns into 2 list of categorical and numerical data features**.

Finding out the existance of null values in dataset

In [11]:

```
number_of_null_values_for_every_column = train_df.isnull().sum()
percentage_of_null_values_for_every_column = ((100*train_df.isnull().sum())/train_df.count())
null_percentage_df = pd.concat([number_of_null_values_for_every_column, percentage_of_null_values_for_every_column], axis=1)
null_percentage_df
```

Out[11]:

	No. of missing Values	Percentage
sku	0	0.000000
national_inv	1	0.000059
lead_time	100894	6.357662
in_transit_qty	1	0.000059
forecast_3_month	1	0.000059
forecast_6_month	1	0.000059
forecast_9_month	1	0.000059
sales_1_month	1	0.000059
sales_3_month	1	0.000059
sales_6_month	1	0.000059
sales_9_month	1	0.000059
min_bank	1	0.000059
potential_issue	1	0.000059
pieces_past_due	1	0.000059
perf_6_month_avg	1	0.000059
perf_12_month_avg	1	0.000059
local_bo_qty	1	0.000059
deck_risk	1	0.000059
oe_constraint	1	0.000059
ppap_risk	1	0.000059
stop_auto_buy	1	0.000059
rev_stop	1	0.000059
went_on_backorder	1	0.000059

- From the above code snippet we can say that the **lead_time column contain 100894 null values out of 1687861 which is 6.357%** of total number of data and it is significant number with respect to other columns.
- Another important thing is noticed that the **dependent variable of type categorical is also include a nan value for a single row**, which we have to remove for better prediction.


Findout the record which have dependent variable as "nan" and remove it from original dataset

```
In [12]: train_df[train_df['went_on_backorder'].isna()]
```

Out[12]:

	sku	national_inv	lead_time	in_transit_qty	forecast_3_month	forecast_6_month	forecast_9_month
1687860	(1687860 rows)	NaN	NaN	NaN	NaN	NaN	NaN

1 rows × 23 columns



```
In [13]: train_df.drop(index = 1687860,inplace = True)
```

```
In [14]: train_df['lead_time'].fillna(train_df['lead_time'].mean(),inplace=True)
```

Cross check the null values

```
In [15]: number_of_null_values_for_every_column = train_df.isnull().sum()
percentage_of_null_values_for_every_column = ((100*train_df.isnull().sum())/train_df.count())
null_percentage_df = pd.concat([number_of_null_values_for_every_column, percentage_of_null_valu
null_percentage_df
```

Out[15]:

	No. of missing Values	Percentage
sku	0	0.0
national_inv	0	0.0
lead_time	0	0.0
in_transit_qty	0	0.0
forecast_3_month	0	0.0
forecast_6_month	0	0.0
forecast_9_month	0	0.0
sales_1_month	0	0.0
sales_3_month	0	0.0
sales_6_month	0	0.0
sales_9_month	0	0.0
min_bank	0	0.0
potential_issue	0	0.0
pieces_past_due	0	0.0
perf_6_month_avg	0	0.0
perf_12_month_avg	0	0.0
local_bo_qty	0	0.0
deck_risk	0	0.0
oe_constraint	0	0.0
ppap_risk	0	0.0

	No. of missing Values	Percentage
stop_auto_buy	0	0.0
rev_stop	0	0.0
went_on_backorder	0	0.0

Basic statistics of features

In [16]: `train_df[numarical_features].describe().transpose()`

	count	mean	std	min	25%	50%	75%	max
national_inv	1687860.0	496.111782	29615.233831	-27256.0	4.00	15.00	80.00	12334404.0
lead_time	1687860.0	7.872267	6.841885	0.0	4.00	8.00	8.00	52.0
in_transit_qty	1687860.0	44.052022	1342.741731	0.0	0.00	0.00	0.00	489408.0
forecast_3_month	1687860.0	178.119284	5026.553102	0.0	0.00	0.00	4.00	1427612.0
forecast_6_month	1687860.0	344.986664	9795.151861	0.0	0.00	0.00	12.00	2461360.0
forecast_9_month	1687860.0	506.364431	14378.923562	0.0	0.00	0.00	20.00	3777304.0
sales_1_month	1687860.0	55.926069	1928.195879	0.0	0.00	0.00	4.00	741774.0
sales_3_month	1687860.0	175.025930	5192.377625	0.0	0.00	1.00	15.00	1105478.0
sales_6_month	1687860.0	341.728839	9613.167104	0.0	0.00	2.00	31.00	2146625.0
sales_9_month	1687860.0	525.269701	14838.613523	0.0	0.00	4.00	47.00	3205172.0
min_bank	1687860.0	52.772303	1254.983089	0.0	0.00	0.00	3.00	313319.0
pieces_past_due	1687860.0	2.043724	236.016500	0.0	0.00	0.00	0.00	146496.0
perf_6_month_avg	1687860.0	-6.872059	26.556357	-99.0	0.63	0.82	0.97	1.0
perf_12_month_avg	1687860.0	-6.437947	25.843331	-99.0	0.66	0.81	0.95	1.0
local_bo_qty	1687860.0	0.626451	33.722242	0.0	0.00	0.00	0.00	12530.0

The above list shows the **basic statistics of numarical features** which includes the **total number of values for each column, minimum, maximum values for each columns, percentile values etc.**

In [17]: `train_df[categorical_features].describe().transpose()`

	count	unique	top	freq
sku	1687860	1687860	1375557	1
potential_issue	1687860	2	No	1686953
deck_risk	1687860	2	No	1300377
oe_constraint	1687860	2	No	1687615
ppap_risk	1687860	2	No	1484026
stop_auto_buy	1687860	2	Yes	1626774
rev_stop	1687860	2	No	1687129
went_on_backorder	1687860	2	No	1676567

The above list shows the **basic statistics of categorical features** which includes the **total number of values for each column, total unique values for each columns, top most values etc.**

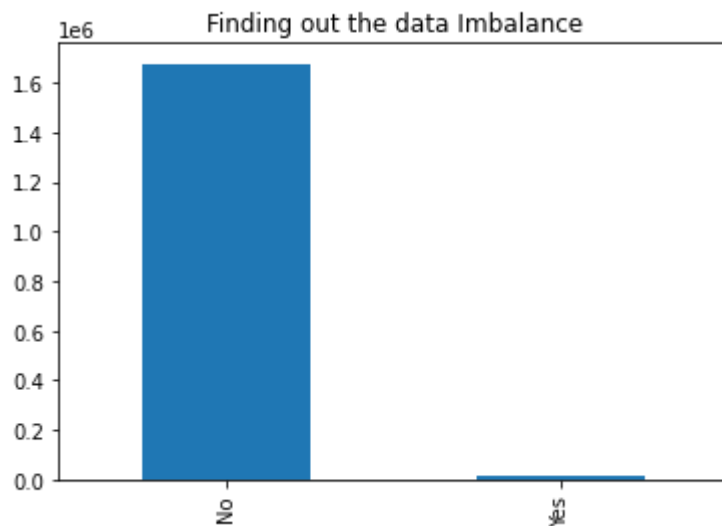
From the above snippet we can say

1. **"lead_time"** column has large number of nan values (missing values) with respect to the other columns of the dataset which is replaced by the mean of that column data.
2. From the **"sku"** column statistics we can say that there are total 1687861 numbers of unique product variants available in the given dataset.
3. Most of the product variants did not went to backorder list as the top value of **went_back_order** feature is **No**.

Specific numbers of product variants records that went backorder

```
In [18]: train_df['went_on_backorder'].value_counts().plot.bar()
          plot.title("Finding out the data Imbalance")
          train_df['went_on_backorder'].value_counts()
```

```
Out[18]: No      1676567
          Yes       11293
          Name: went_on_backorder, dtype: int64
```



Total no of **training data points are 1687860** and among them **11293 data had gone on back order**. So it implies that only **0.66%** of the training dataset **has gone on backorder**, which is **highly imbalanced**.

Univariate analysis

```
In [19]: for feature in categorical_features:
          print(feature+" has following uinique values:")
          print(train_df[feature].unique())
          print("-"*100)
```

```
sku has following uinique values:
[1026827 1043384 1043696 ... '1524346' '1439563' '1502009']
```

```
-----
potential_issue has following uinique values:
['No' 'Yes']
```

```
-----
deck_risk has following uinique values:
['No' 'Yes']
```

```
-----
oe_constraint has following uinique values:
['No' 'Yes']
```

```
-----
ppap_risk has following uinique values:
```

```
['No' 'Yes']
```

```
stop_auto_buy has following unique values:  
['Yes' 'No']
```

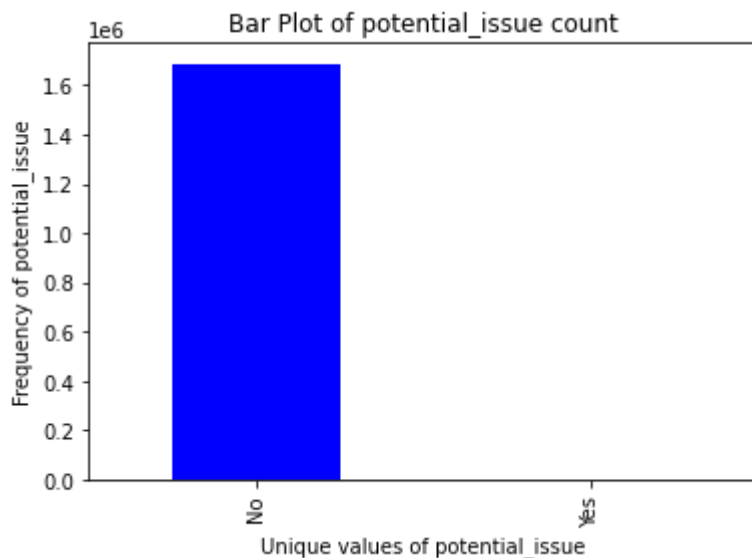
```
rev_stop has following unique values:  
['No' 'Yes']
```

```
went_on_backorder has following unique values:  
['No' 'Yes']
```

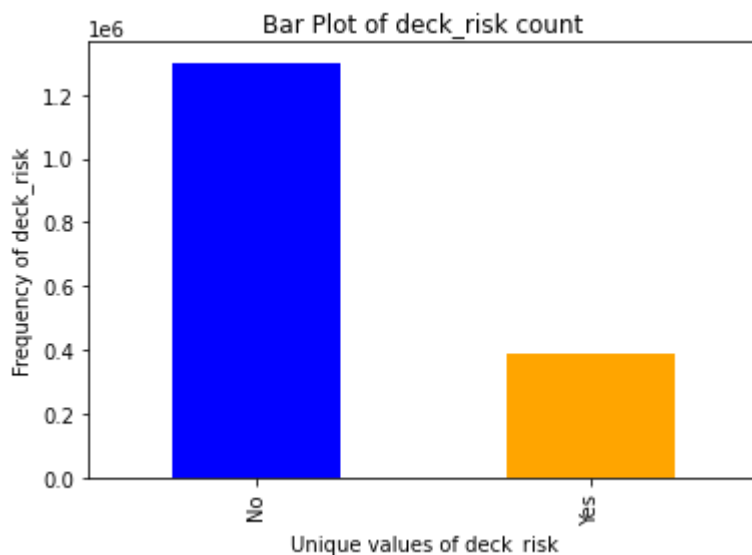
We can see that except "sku" all the **categorical variables** has only **2 unique values Yes and No**

In [20]:

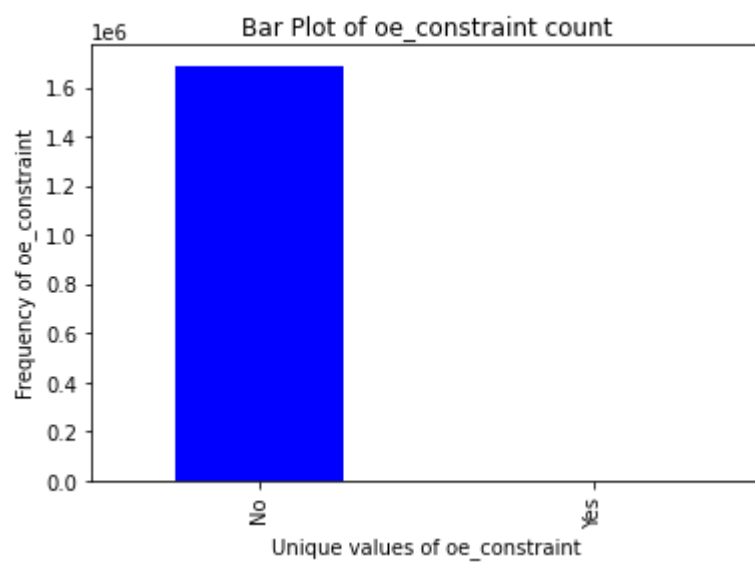
```
for feature in categorical_features:  
    if(feature != 'sku'):  
        train_df[feature].value_counts().plot.bar(color=['blue','orange'])  
        plot.title("Bar Plot of "+feature+" count")  
        plot.ylabel("Frequency of "+feature)  
        plot.xlabel("Unique values of "+ feature)  
        plot.show()  
        print(train_df[feature].value_counts())  
        plot.close()
```



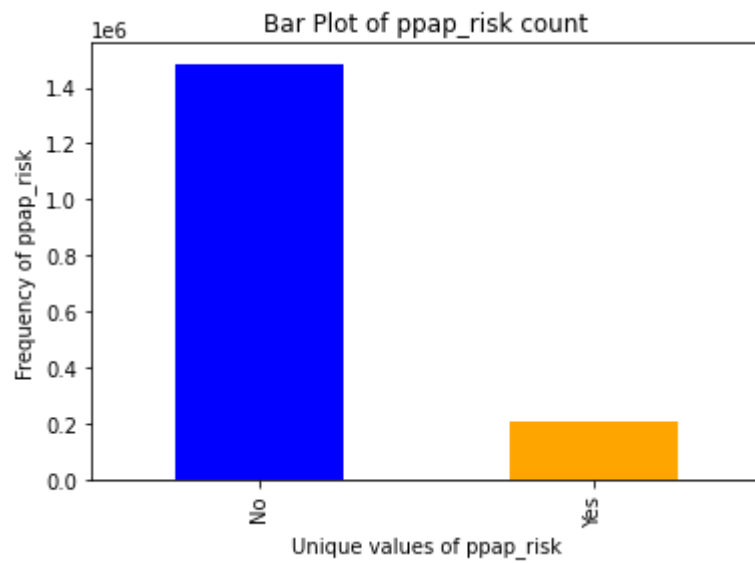
```
No      1686953  
Yes       907  
Name: potential_issue, dtype: int64
```



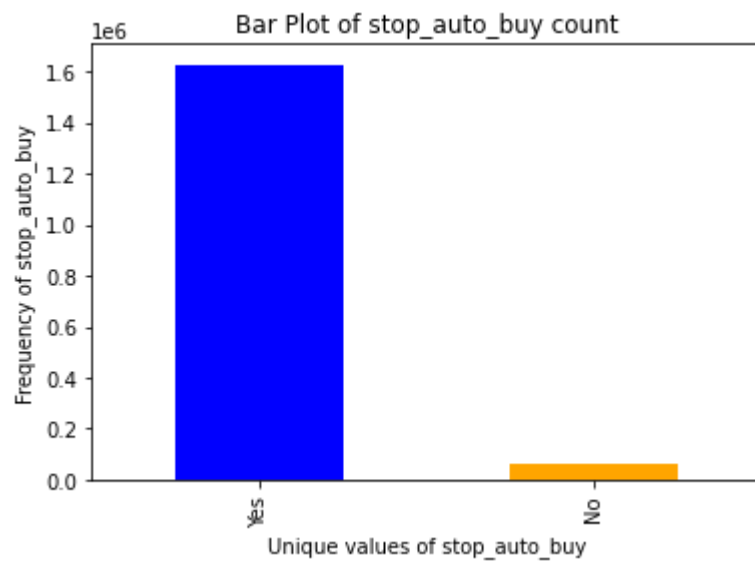
```
No      1300377  
Yes       387483  
Name: deck_risk, dtype: int64
```



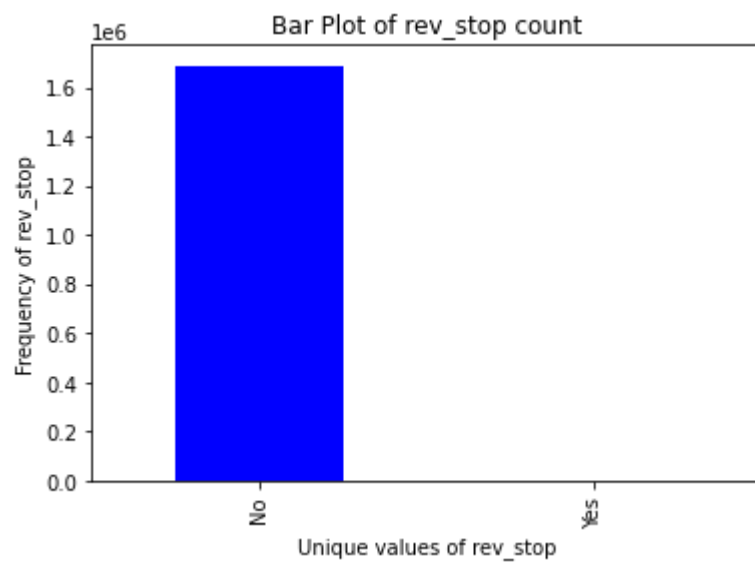
No 1687615
 Yes 245
 Name: oe_constraint, dtype: int64



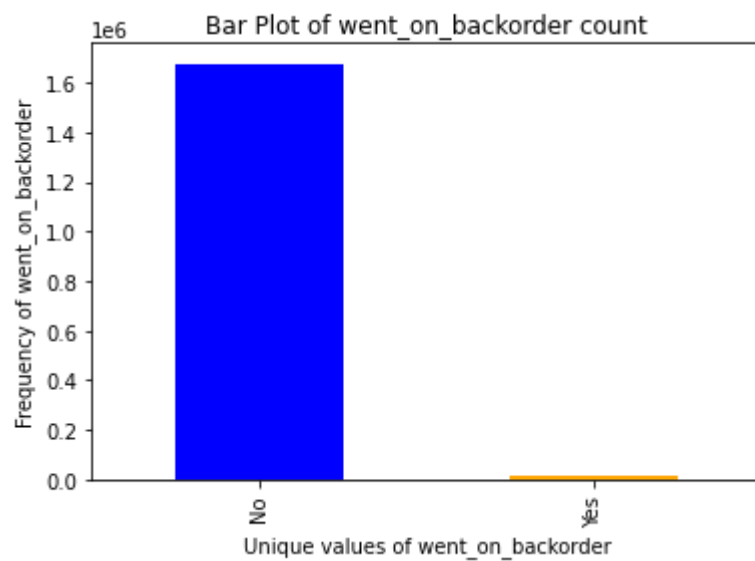
No 1484026
 Yes 203834
 Name: ppap_risk, dtype: int64



Yes 1626774
 No 61086
 Name: stop_auto_buy, dtype: int64




No 1687129
 Yes 731
 Name: rev_stop, dtype: int64



No 1676567
 Yes 11293
 Name: went_on_backorder, dtype: int64

For all the categorical variables No is majority. Also in context to all the categorical variable data is highly imbalance.

```
In [21]: # sns.kdeplot(train_df['lead_time'], hue=train_df['went_on_backorder'])
# plot.title("Histogram Plot of Lead time")
# plot.show()
# plot.close()
sns.boxplot(x='went_on_backorder',y='lead_time', data=train_df).set_title('Box plot of lead_time')
plot.show()
plot.close()
```

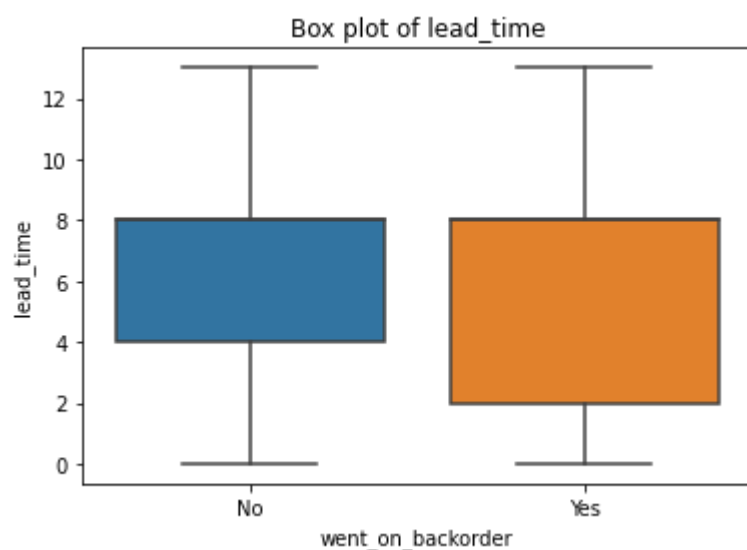


- On the top whisker of the majority class there are few outliers which need to be removed.
- The 50th percentile is almost similar to the 75th percentile as we replace the missing value with the mean.

```
In [22]: def remove_outlier(df_in, col_name):
          q1 = df_in[col_name].quantile(0.25)
          q3 = df_in[col_name].quantile(0.75)
          iqr = q3-q1 #Interquartile range
          fence_low = q1-1.5*iqr
          fence_high = q3+1.5*iqr
          df_out = df_in.loc[(df_in[col_name] > fence_low) & (df_in[col_name] < fence_high)]
          return df_out
```

```
In [23]: train_df = remove_outlier(train_df, 'lead_time')
```

```
In [24]: sns.boxplot(x='went_on_backorder',y='lead_time', data=train_df).set_title('Box plot of lead_time')
          plot.show()
          plot.close()
```



```
In [25]: train_df.shape
```

```
Out[25]: (1629448, 23)
```

So after removing the outlier the number of rows of the dataset is reduced down to **1629448** and from the above boxplot we can see that there is no such outliers.

Find out correlation among the features


In [26]:

```
corr_matrix = train_df[numarical_features].corr()

mask = np.zeros_like(corr_matrix, dtype=np.bool)
mask[np.triu_indices_from(mask)] = True

plot.figure(figsize = (15,15))
plot.title('Correlation Matrix')
sns.heatmap(corr_matrix,
            vmin=-1,
            cmap='viridis',
            annot=True,
            mask = mask)

plot.show()
```



From the above correlation matrix we can observe that the following independent variable have high correlation, which are as follows -


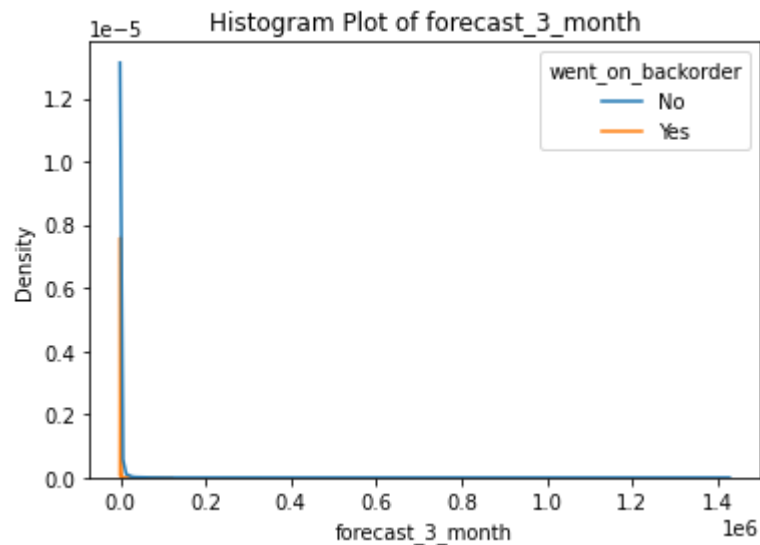
- **forecast_3_month, forecast_6_month, forecast_9_month** have correlation at around 0.99. So we can say those 3 features will show nearly same behaviour. We can pick one of them for the final model.

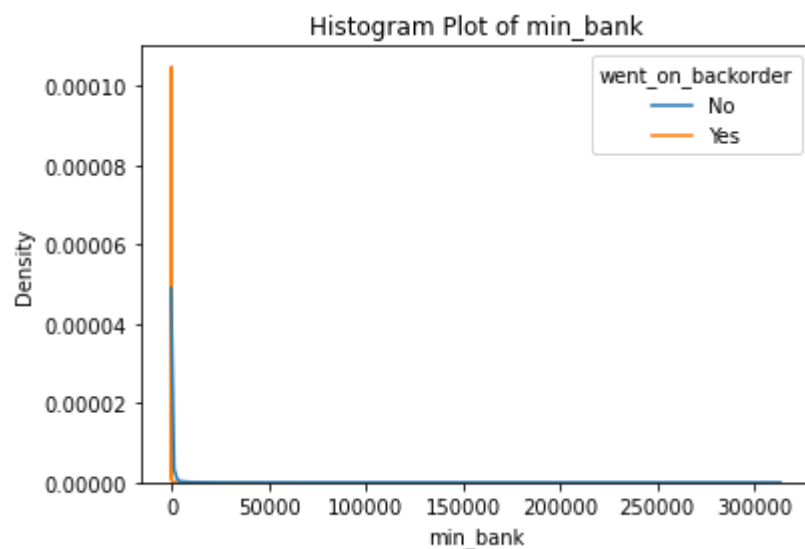
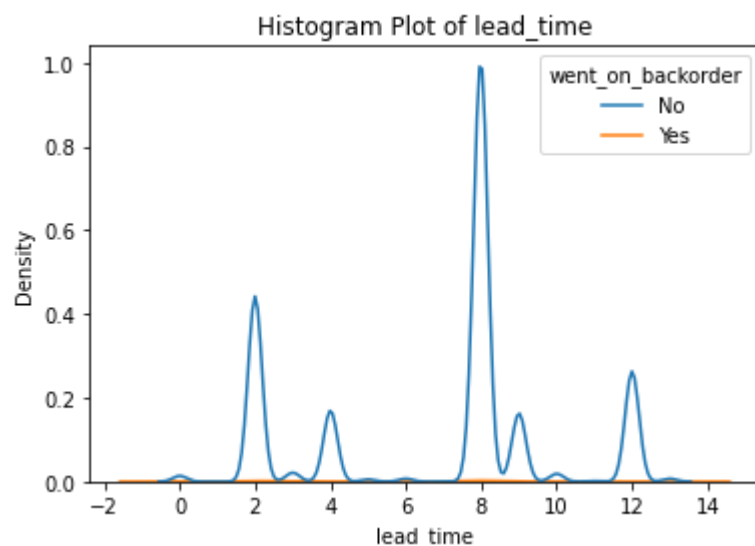

- **sales_1_month, sales_3_month, sales_6_month, sales_9_month** have correlation ranges from 0.82 to 0.92. So we can say those 4 features will show nearly same behaviour. We can pick one of them for the final model.
- **pref_6_month_avg, pref_12_month_avg** have correlation at around 0.97. So we can say those 2 features will show nearly same behaviour. We can pick one of them for the final model.

Impact of numerical features with respect to target

In [27]:

```
numerical_features_subset = ['forecast_3_month', 'sales_1_month', 'perf_6_month_avg', 'lead_time']
for feature in numerical_features_subset:
    sns.kdeplot(train_df[feature], hue=train_df['went_on_backorder'])
    plot.title("Histogram Plot of "+feature)
    plot.show()
    plot.close()
```





From the above plots we find that when the prediction for the future sale is low as well as the past sale quantity is very low then only the possibility of backorder happen. Which implied that products which have low sales and low demand there is a possibility of went on backorder.

Impact of categorical features with respect to the target

In [28]:

```
for feature in categorical_features:
    if(feature != 'sku' and feature != 'went_on_backorder'):
        counts_df = train_df.groupby([feature, "went_on_backorder"])[ "sku"].count().unstack()
        percentage_df = counts_df.T.div(counts_df.T.sum()).T
        print(percentage_df)
        fig, ax = plot.subplots()
        percentage_df.plot(kind="bar", stacked=True, ax=ax)
```

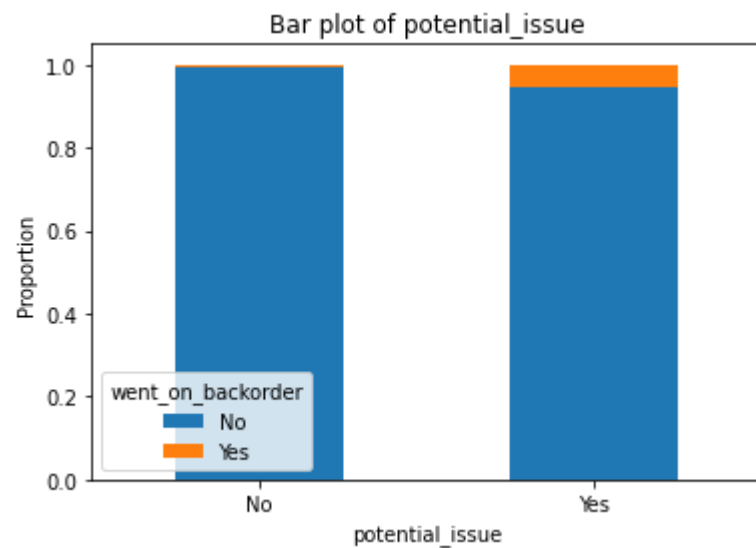


```

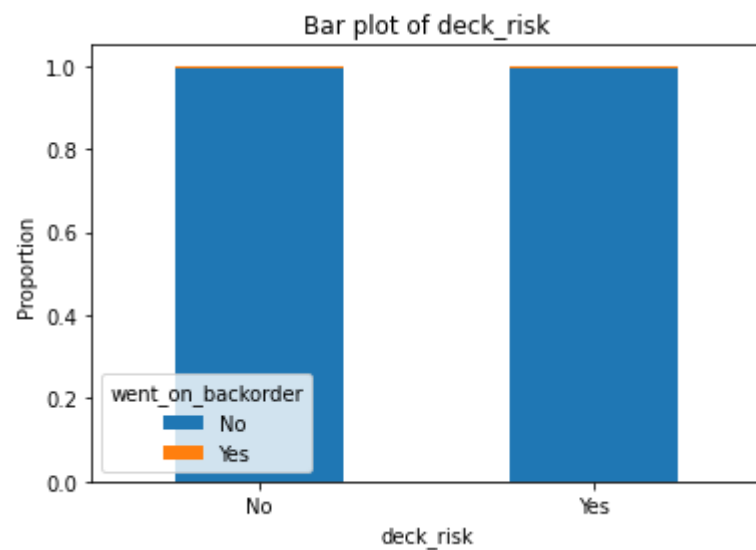
ax.set_xlabel(feature)
ax.set_xticklabels(['No', 'Yes'], rotation=0)
ax.set_ylabel("Proportion")
ax.set_title("Bar plot of "+feature)
plot.show()
plot.close()

```

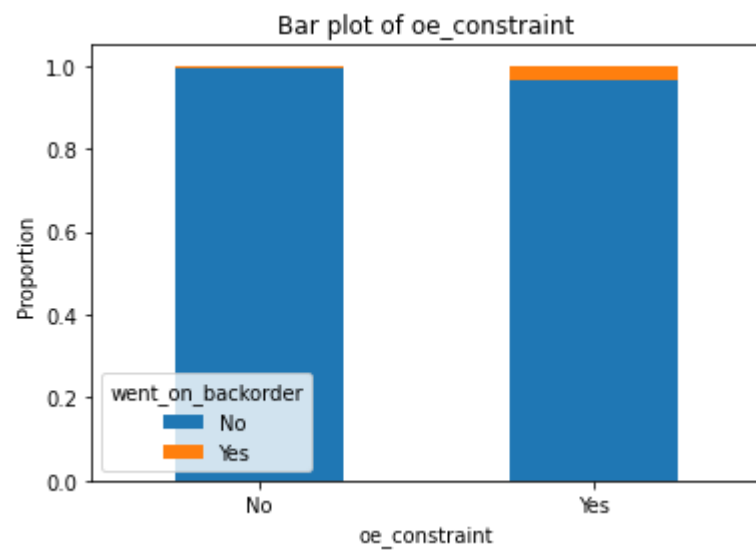
went_on_backorder	No	Yes
potential_issue		
No	0.993189	0.006811
Yes	0.943630	0.056370



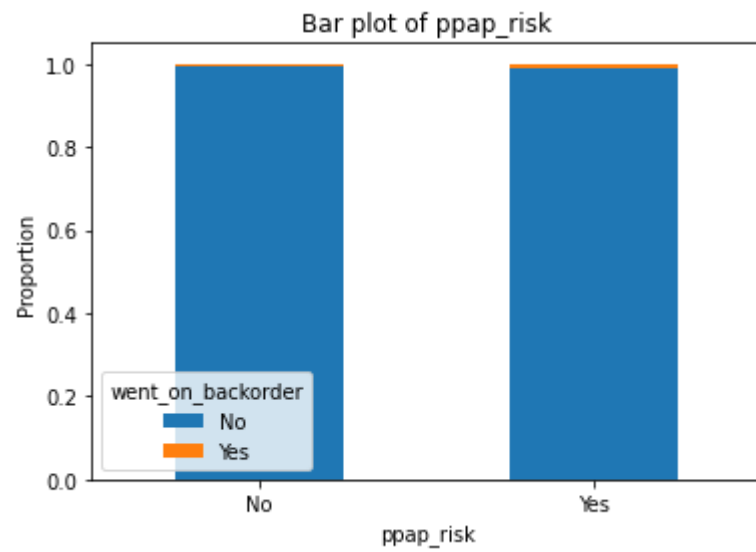
went_on_backorder	No	Yes
deck_risk		
No	0.992729	0.007271
Yes	0.994706	0.005294



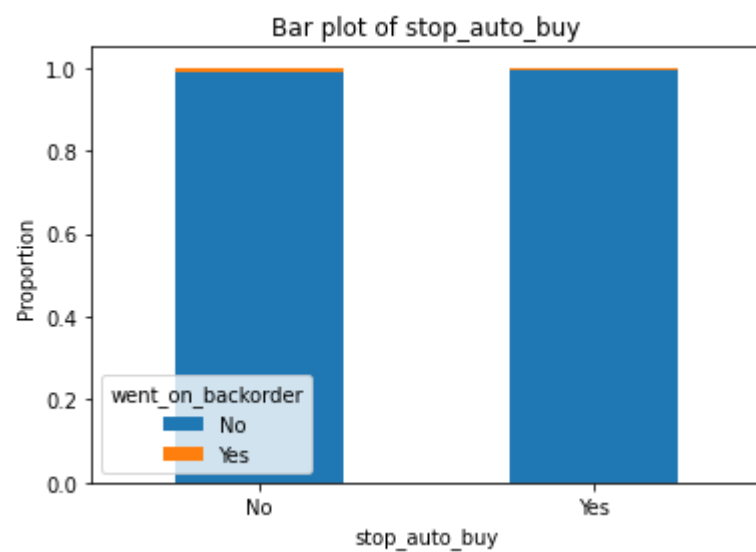
went_on_backorder	No	Yes
oe_constraint		
No	0.993166	0.006834
Yes	0.967347	0.032653



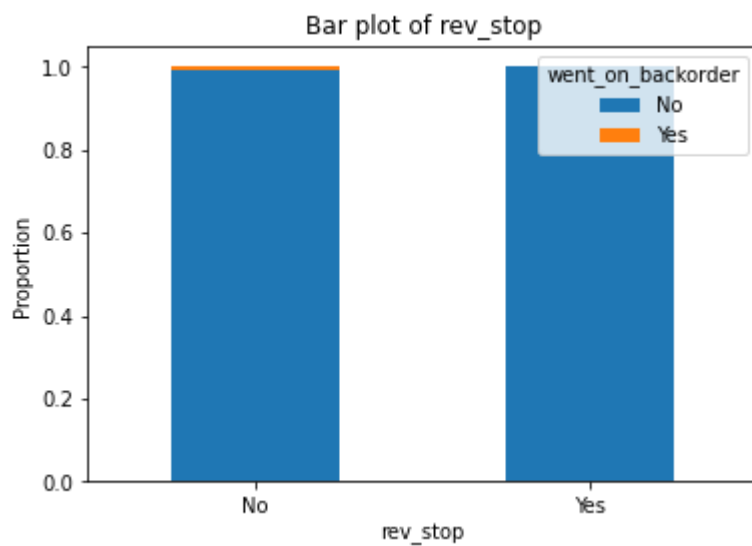
went_on_backorder	No	Yes
ppap_risk		
No	0.993446	0.006554
Yes	0.991076	0.008924



went_on_backorder	No	Yes
stop_auto_buy		
No	0.990632	0.009368
Yes	0.993238	0.006762



went_on_backorder	No	Yes
rev_stop		
No	0.993159	0.006841
Yes	1.000000	NaN

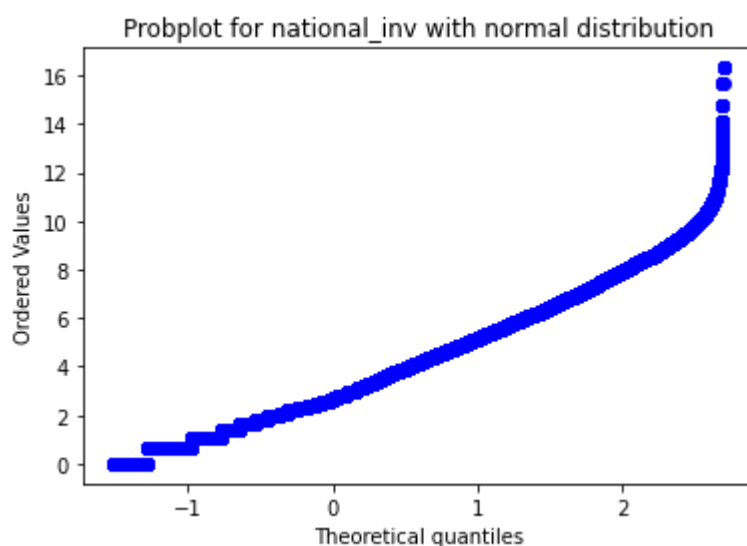



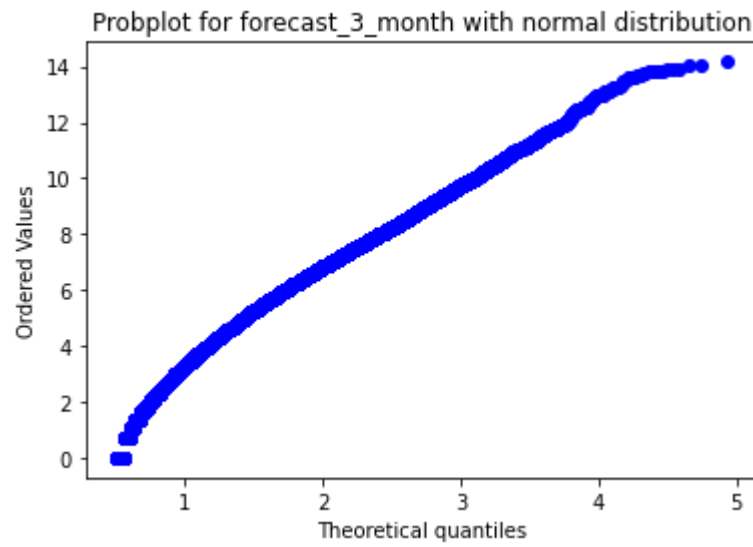
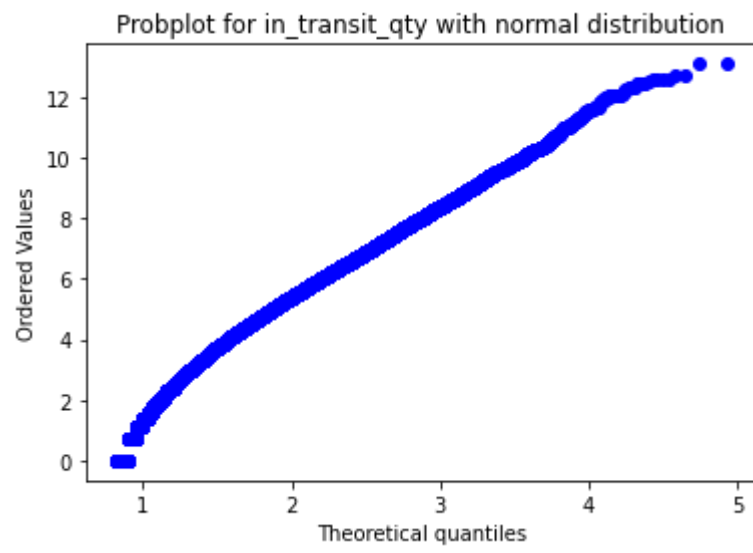
- If potential_issue flag is 'Yes' then there is 5.6% chance of item going to back order.
- For the changes in deck_risk flag there is no such significant changes noted in the decision making whether the product will went on back order or not.
- If potential_issue flag is 'Yes' then there is 3.26% chance of item going to back order.
- For the changes in ppap_risk flag there is no such significant changes noted in the decision making whether the product will went on back order or not.
- For the changes in stop_auto_buy flag there is no such significant changes noted in the decision making whether the product will went on back order or not.
- For rev_stop no item went to backorder if rev_stop flag is set to Yes. For rev_stop = 'No' the proportion of orders that did go to back order and that weren't is same as class ratio that is 99.33 % : 0.667%

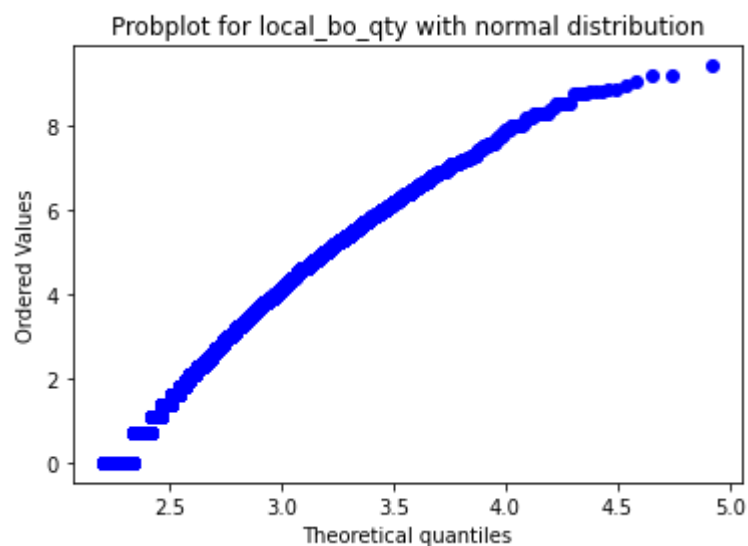
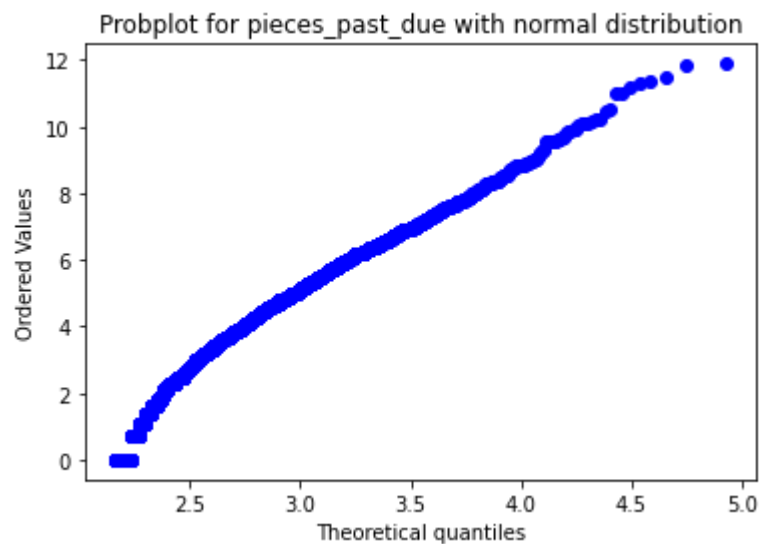
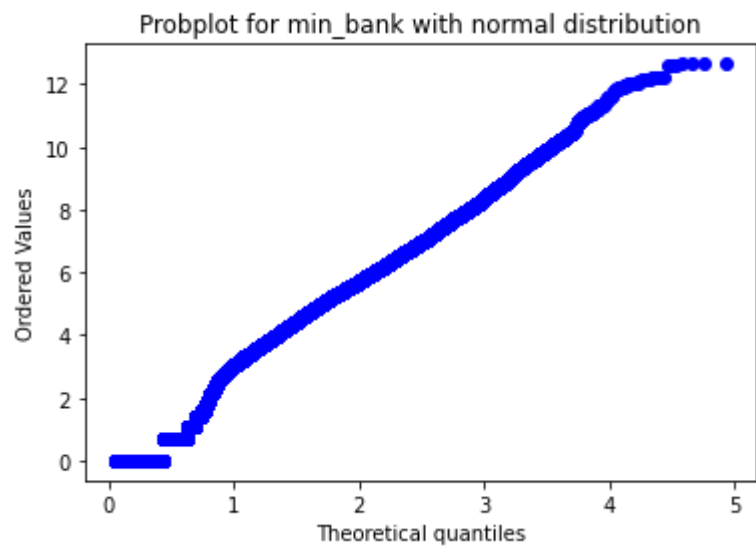
As the numarical features are right skewed so we can check whether they follow log normal distribution or not by using QQ plot

In [29]:

```
skewed_right_subset = ['national_inv', 'in_transit_qty', 'forecast_3_month', 'sales_3_month',
for feature in skewed_right_subset:
    fgr = plot.figure()
    ax = fgr.add_subplot(111)
    sa = stats.probplot(np.log(train_df[feature].values),plot=ax)
    ax.set_title("Probplot for "+feature+" with normal distribution")
    plot.show()
```







As we got approximately straight line with 45 degree angle on the plot between log transformed feature and the normal distribution we can say that the above subset of features are fall under log normal distribution.

Feature Engineering

In performance feature columns there are some values which is (-99.0) unexpected for that feature. We can say that it is some missing values and we replaced it by NaN

```
In [30]: train_df.perf_12_month_avg.replace({-99.0 : np.nan}, inplace = True)
          train_df.perf_6_month_avg.replace({-99.0 : np.nan}, inplace = True)
```

```
In [31]: for feature in categorical_features:
```

```
if(feature != 'sku'):
    train_df[feature] = train_df[feature].map({"Yes" : 1, "No" : 0})
```

Train test split

```
In [32]: y = y_hyp = train_df['went_on_backorder']
X = X_hyp = train_df.drop(['sku', 'went_on_backorder'], axis=1)
```

```
In [33]: X_train, X_test, y_train, y_test = train_test_split(X, y, random_state = 42, stratify=y, test_s
X_train, X_cv, y_train, y_cv = train_test_split(X_train, y_train, random_state = 42, stratify=y,
```

```
In [34]: print(X_train.shape)
# print(X_train_hyp.shape)
print(X_cv.shape)
print(X_test.shape)
```

```
(1173202, 21)
(130356, 21)
(325890, 21)
```

```
In [35]: print(y_train.shape)
# print(y_train_hyp.shape)
print(y_cv.shape)
print(y_test.shape)
```

```
(1173202,)
(130356,)
(325890,)
```

Findout the column index which have right skewed data

```
In [36]: skewed_right = ['national_inv', 'in_transit_qty', 'forecast_3_month', 'forecast_6_month', 'fore
log_columns_index = []
for i in skewed_right:
    log_columns_index.append(X_train.columns.get_loc(i))
```

Impute the missing value with Iterative Imputer

This imputer is an advanced version of standard imputer as the mechanism it used to impute the missing values is as follows -

It will find all the rows of a dataframe which donot have a particular feature missing and make a regression model on top of it and finally predict the missing value.

```
In [37]: # Ref - https://scikit-learn.org/stable/modules/generated/sklearn.impute.IterativeImputer.html
iterativeImputr = IterativeImputer()
iterativeImputr.fit(X_train)
X_train = iterativeImputr.transform(X_train)
# X_train_hyp = iterativeImputr.transform(X_train_hyp)
X_test = iterativeImputr.transform(X_test)
X_cv = iterativeImputr.transform(X_cv)

X_train_log, X_cv_log, X_test_log, y_train_log, y_cv_log, y_test_log = X_train, X_cv, X_test, y
```

Apply log transform on right skewed train data

```
In [38]:
```

#Ref - <https://towardsdatascience.com/transforming-skewed-data-73da4c2d0d16>

```
def log_transform(a):  
    sign = np.sign(a[log_columns_index])  
    a[log_columns_index] = np.log(1.0+abs(a[log_columns_index]))*sign  
    return a
```

```
In [39]: X_train_log = np.apply_along_axis(log_transform, 1, X_train_log)  
X_cv_log = np.apply_along_axis(log_transform, 1, X_cv_log)  
X_test_log = np.apply_along_axis(log_transform, 1, X_test_log)
```

Apply robust scaler to scale data having outlire

Robust scaler is used as because it can scale the data by avoiding the impact of outlire data using the formulae value = (value – median) / (p75 – p25)

```
In [40]: robusrscaler = RobustScaler()  
robusrscaler.fit(X_train)  
X_train = robusrscaler.transform(X_train)  
# X_train_hyp = robusrscaler.transform(X_train_hyp)  
X_cv = robusrscaler.transform(X_cv)  
X_test = robusrscaler.transform(X_test)
```

Apply standard scalar on log transformed data

```
In [41]: standardscalar = StandardScaler()  
standardscalar.fit(X_train_log)  
X_train_log = standardscalar.transform(X_train_log)  
X_cv_log = standardscalar.transform(X_cv_log)  
X_test_log = standardscalar.transform(X_test_log)
```

Apply Different Models

Accuracy Metric Check

```
In [42]: def plot_confusion_matrix(test_y, predict_y, dataset):  
    C = confusion_matrix(test_y, predict_y)  
    labels = [0,1]  
    cmap=sns.light_palette("green")  
    # representing A in heatmap format  
    print("-"*40, dataset+" Confusion matrix", "-"*40)  
    plot.figure(figsize=(3,3))  
    sns.heatmap(C, annot=True, cmap=cmap, fmt=".3f", xticklabels=labels, yticklabels=labels)  
    plot.xlabel('Predicted Class')  
    plot.ylabel('Original Class')  
    plot.show()
```

```
In [43]: #Ref - https://machinelearningmastery.com/roc-curves-and-precision-recall-curves-for-classification  
def accuracy_check(classifier, x_true, y_true):  
  
    plot_confusion_matrix(y_true, classifier.predict(x_true), 'Test')  
    threshold = [0,0.05,0.1,0.15,0.2,0.25,0.3,0.35,0.4,0.45,0.5,0.55,0.6,0.65,0.7,0.75,0.8,0.85,0.9,0.95,1]  
    y_pred = classifier.predict_proba(x_true)[:,-1]  
  
    pr_rc_scores = []  
    tpr = []  
    fpr = []  
    #Check for every threshold  
    for thr in threshold:  
        pred_classes =[]  
        #for every predictions
```

```

for pred in y_pred:
    if pred > thr:
        pred_classes.append(1)
    else:
        pred_classes.append(0)
pr = precision_score(y_true, pred_classes, pos_label=1)
rc = recall_score(y_true, pred_classes, pos_label=1)
pr_rc_scores.append([rc, pr])

#Calculate FP, TP, FN, TN
pred_classes = np.array(pred_classes)
fp = np.sum((pred_classes == 1) & (y_true == 0))
tp = np.sum((pred_classes == 1) & (y_true == 1))
fn = np.sum((pred_classes == 0) & (y_true == 1))
tn = np.sum((pred_classes == 0) & (y_true == 0))

#Calculate FPR, TPR
fpr.append(fp / (fp + tn))
tpr.append(tp / (tp + fn))

#Calculate F1 score
f1 = f1_score(y_true, classifier.predict(x_true), average = None)
print("The F1 scores : " , f1)
print("The Macro F1 score : ", (f1[0]+f1[1])/2)

#Plot Precision Recall AUC Curve
recall_scores = [X[0] for X in pr_rc_scores]
precision_scores = [Y[1] for Y in pr_rc_scores]
fig = plot.figure(figsize=(8,10))

ax1 = fig.add_subplot(311)
ax1.plot(recall_scores, precision_scores, label = 'PR-AUC Curve'+str(np.round(auc(recall_scores, precision_scores), 3)))
ax1.set_title("Precision - Recall curve")
ax1.set_xlabel("Recall")
ax1.set_ylabel("Precision")
ax1.legend()
plot.subplots_adjust(left=None, bottom=None, right=None, top=None, wspace=0.2, hspace=0.9)

#Plot ROC AUC Curve
ax2 = fig.add_subplot(312)
ax2.plot(fpr, tpr, label = "ROC AUC Curve : "+str(np.round(auc(fpr, tpr), 3)))
ax2.plot([0, 1], ls="--", label='No Skill')
ax2.set_title("ROC-AUC curve")
ax2.set_xlabel("FPR")
ax2.set_ylabel("TPR")
ax2.legend()
plot.subplots_adjust(left=None, bottom=None, right=None, top=None, wspace=0.2, hspace=0.9)

```

Decision Tree Classifier on Robust scaled Data

In [42]:

```

%%time
model = DecisionTreeClassifier()
parameters = {
    'criterion' : ['gini', 'entropy'],
    'max_depth' : [2,4,6,8,10,12]
}
clf = RandomizedSearchCV(model, parameters, cv = 3, verbose = 10, scoring = 'roc_auc', n_iter=100)
gridsrch = clf.fit(X_train, y_train)
print("Best Params : " , gridsrch.best_params_)
print("Best Score : " , gridsrch.best_score_)

```

Fitting 3 folds for each of 5 candidates, totalling 15 fits
 Best Params : {'max_depth': 8, 'criterion': 'gini'}
 Best Score : 0.9238809824762687
 CPU times: user 8.66 s, sys: 368 ms, total: 9.02 s
 Wall time: 22.7 s


In [43]:


```

%%time
model = DecisionTreeClassifier(
    criterion = gridsrch.best_params_['criterion'],
    max_depth = gridsrch.best_params_['max_depth'])
model.fit(X_train, y_train)
plot_confusion_matrix(y_train, model.predict(X_train), 'Training')
print("The ROC-AUC Score obtained on CV set is : ", roc_auc_score(y_cv, model.predict(X_cv)))
print("The F1 scores of each class on CV set are : ", f1_score(y_cv, model.predict(X_cv), average='macro'))
print("The Macro F1-Score obtained on CV set is : ", f1_score(y_cv, model.predict(X_cv), average='macro'))
plot_confusion_matrix(y_cv, model.predict(X_cv), 'Cross Validation')
accuracy_check(model, X_test, y_test)

```

----- Training Confusion matrix -----




The ROC-AUC Score obtained on CV set is : 0.5099890107189377


The F1 scores of each class on CV set are : [0.99652775 0.03837953]

The Macro F1-Score obtained on CV set is : 0.5174536409809918

----- Cross Validation Confusion matrix -----



----- Test Confusion matrix -----





The F1 scores : [0.99656019 0.03955288]

The Macro F1 score : 0.518056533981037

CPU times: user 22.9 s, sys: 61.3 ms, total: 22.9 s

Wall time: 22.9 s




Random Forest Classifier on Robust scaled Data

```
In [44]: %%time
model = RandomForestClassifier(class_weight = "balanced_subsample" , n_jobs = -1)
parameters = {
    'n_estimators' : [10,50,100,300,500,1000],
    'max_depth' : [1,3,5,7]
}
clf = RandomizedSearchCV(model, parameters, cv = 3, verbose = 10, scoring = 'roc_auc', n_iter=100)
gridsrch = clf.fit(X_train, y_train)
print("Best Params : " , gridsrch.best_params_)
print("Best Score : " , gridsrch.best_score_)
```

Fitting 3 folds for each of 5 candidates, totalling 15 fits
 Best Params : {'n_estimators': 100, 'max_depth': 7}
 Best Score : 0.9361650843187165
 CPU times: user 3min 28s, sys: 2.24 s, total: 3min 30s
 Wall time: 14min 10s

```
In [45]: %%time
model = RandomForestClassifier(
    n_estimators = gridsrch.best_params_['n_estimators'],
    max_depth = gridsrch.best_params_['max_depth'],
    class_weight = "balanced_subsample",
    n_jobs = -1)
model.fit(X_train, y_train)
plot_confusion_matrix(y_train, model.predict(X_train), 'Training')
print("The ROC-AUC Score obtained on CV set is : " , roc_auc_score(y_cv, model.predict(X_cv)))
print("The F1 scores of each class on CV set are : " , f1_score(y_cv, model.predict(X_cv), average='macro'))
print("The Macro F1-Score obtained on CV set is : " , f1_score(y_cv, model.predict(X_cv), average='macro'))
plot_confusion_matrix(y_cv, model.predict(X_cv), 'Cross Validation')
accuracy_check(model, X_test, y_test)
```

----- Training Confusion matrix -----




The ROC-AUC Score obtained on CV set is : 0.8675173704370785


The F1 scores of each class on CV set are : [0.92214269 0.0771999]

The Macro F1-Score obtained on CV set is : 0.499671294603133

----- Cross Validation Confusion matrix -----



----- Test Confusion matrix -----





The F1 scores : [0.92100123 0.07484837]

The Macro F1 score : 0.4979247999345704

CPU times: user 4min 24s, sys: 1.38 s, total: 4min 25s

Wall time: 1min 1s



Balanced Bagging classifier on Robust scaled Data

```
In [46]: %%time
model = BalancedBaggingClassifier(n_jobs = -1)
parameters = {'n_estimators' : [10,50,100,300,500,1000]}
clf = RandomizedSearchCV(model, parameters, cv = 3, verbose = 10, scoring = 'roc_auc', n_iter=100)
gridsrch = clf.fit(X_train, y_train)
print("Best Params : " , gridsrch.best_params_)
print("Best Score : " , gridsrch.best_score_)
```

Fitting 3 folds for each of 5 candidates, totalling 15 fits

Best Params : {'n_estimators': 1000}

Best Score : 0.9650438438505589


CPU times: user 972 ms, sys: 1.05 s, total: 2.02 s

Wall time: 18min 10s

```
In [47]: %%time
model = BalancedBaggingClassifier(
    n_estimators = gridsrch.best_params_['n_estimators'],
    n_jobs = -1)
model.fit(X_train, y_train)
plot_confusion_matrix(y_train, model.predict(X_train), 'Training')
print("The ROC-AUC Score obtained on CV set is : " , roc_auc_score(y_cv, model.predict(X_cv)))
print("The F1 scores of each class on CV set are : " , f1_score(y_cv, model.predict(X_cv), average='macro'))
print("The Macro F1-Score obtained on CV set is : " , f1_score(y_cv, model.predict(X_cv), average='macro'))
plot_confusion_matrix(y_cv, model.predict(X_cv), 'Cross Validation')

accuracy_check(model,X_test,y_test)
```

----- Training Confusion matrix -----




The ROC-AUC Score obtained on CV set is : 0.9199787236283586


The F1 scores of each class on CV set are : [0.95665411 0.13250584]

The Macro F1-Score obtained on CV set is : 0.5445799791872363

----- Cross Validation Confusion matrix -----



----- Test Confusion matrix -----





The F1 scores : [0.95623685 0.13123041]

The Macro F1 score : 0.5437336253843814

CPU times: user 25.4 s, sys: 2.39 s, total: 27.8 s

Wall time: 7min 20s



XGBoost classifier on Robust scaled Data

In [48]: `!pip3 install tqdm`

Requirement already satisfied: tqdm in /opt/conda/lib/python3.7/site-packages (4.60.0)

```
In [49]: import random
from tqdm import tqdm
def custom_loops(x_train,y_train,classifier, param_range, folds):
    #    trainscores = []
    testscores = []
    #Referance - https://www.geeksforgeeks.org/python-random-sample-function/
    #1.generate 10 unique values(uniform random distribution) in the given range "param_range"

    lst = random.sample(param_range,5)
    lst.sort()

    params = {'n_estimators':lst}

    for k in tqdm(params['n_estimators']):
        testscores_folds = []
        for fold in range(0, folds):
            #2.devide numbers ranging from 0 to len(X_train) into groups= folds
            block_size = int(len(x_train)/folds)
            test_indices = list(set(list(range((block_size*fold), (block_size*(fold+1))))))
            train_indices = list(set(list(range(1, len(x_train)))) - set(test_indices))
            # selecting the data points based on the train_indices and test_indices
            X_train = pd.DataFrame(x_train).iloc[train_indices]
            Y_train = pd.DataFrame(y_train).iloc[train_indices]
            X_test = pd.DataFrame(x_train).iloc[test_indices]
            Y_test = pd.DataFrame(y_train).iloc[test_indices]

            classifier.n_estimators = k
            classifier.fit(X_train,Y_train)

            Y_predicted = classifier.predict(X_test)
            testscores_folds.append(f1_score(Y_test, Y_predicted, average = 'macro'))

        testscores.append(np.mean(np.array(testscores_folds)))
    return testscores,params
```

In [50]:

```
%time
model = XGBClassifier(nthread=-1)
parameters = [10,50,100,300,500,1000]
#clf = RandomizedSearchCV(model, parameters, scoring = 'roc_auc', n_jobs=-1)
#gridsrch = clf.fit(X_train, y_train)

testscores,paramsopt = custom_loops(X_train, y_train, model, parameters, 3)
print("Params : " , paramsopt)
print("Scores : " , testscores)
```

```
0%|          | 0/5 [00:00<?, ?it/s]
[14:42:14] WARNING: ../src/learner.cc:1095: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.
[14:42:18] WARNING: ../src/learner.cc:1095: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.
[14:42:22] WARNING: ../src/learner.cc:1095: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.
20%|██        | 1/5 [00:11<00:46, 11.74s/it]
[14:42:25] WARNING: ../src/learner.cc:1095: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.
[14:42:39] WARNING: ../src/learner.cc:1095: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.
[14:42:53] WARNING: ../src/learner.cc:1095: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.
40%|████      | 2/5 [00:52<01:26, 28.97s/it]
[14:43:06] WARNING: ../src/learner.cc:1095: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.
[14:44:24] WARNING: ../src/learner.cc:1095: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.
[14:45:41] WARNING: ../src/learner.cc:1095: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.
60%|██████    | 3/5 [04:45<04:04, 122.02s/it]
[14:46:59] WARNING: ../src/learner.cc:1095: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.
[14:49:07] WARNING: ../src/learner.cc:1095: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.
[14:51:14] WARNING: ../src/learner.cc:1095: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.
80%|████████  | 4/5 [11:08<03:45, 225.20s/it]
[14:53:23] WARNING: ../src/learner.cc:1095: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.
[14:57:34] WARNING: ../src/learner.cc:1095: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.
[15:01:34] WARNING: ../src/learner.cc:1095: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.
100%|██████████| 5/5 [23:32<00:00, 282.54s/it]
Params : {'n_estimators': [10, 50, 300, 500, 1000]}
Scores : [0.5010244184535769, 0.5461729764736679, 0.6273638740067524, 0.6585902116956097, 0.6978457280301775]
CPU times: user 3h 6min 35s, sys: 1.78 s, total: 3h 6min 36s
Wall time: 23min 32s
```

In [51]:

```
%time
model = XGBClassifier(
    n_estimators = 1000,
```


```

        nthread = -1)
model.fit(X_train, y_train)
plot_confusion_matrix(y_train, model.predict(X_train), 'Training')
print("The ROC-AUC Score obtained on CV set is : " ,roc_auc_score(y_cv, model.predict(X_cv)))
print("The F1 scores of each class on CV set are : ",f1_score(y_cv, model.predict(X_cv),average
print("The Macro F1-Score obtained on CV set is : " , f1_score(y_cv, model.predict(X_cv),average
plot_confusion_matrix(y_cv, model.predict(X_cv), 'Cross Validation')
accuracy_check(model,X_test,y_test)

```

[15:07:45] WARNING: ../src/learner.cc:1095: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

----- Training Confusion matrix -----




The ROC-AUC Score obtained on CV set is : 0.6754097574048912


The F1 scores of each class on CV set are : [0.99753707 0.49486166]

The Macro F1-Score obtained on CV set is : 0.7461993646535316

----- Cross Validation Confusion matrix -----



----- Test Confusion matrix -----





The F1 scores : [0.9974624 0.47512755]

The Macro F1 score : 0.7362949747504652

CPU times: user 51min 54s, sys: 396 ms, total: 51min 54s

Wall time: 6min 51s



Adaboost classifier on Robust scaled Data

In [42]:


```
%%time
model = AdaBoostClassifier()
parameters = {'n_estimators' : [10,50,100,300,500,1000]}
clf = RandomizedSearchCV(model, parameters,scoring = 'roc_auc', n_jobs=-1)
gridsrch = clf.fit(X_train, y_train)
print("Best Params : " , gridsrch.best_params_)
print("Best Score : " , gridsrch.best_score_)
```

Best Params : {'n_estimators': 1000}
 Best Score : 0.9401228188639209
 CPU times: user 20min 54s, sys: 37.6 s, total: 21min 32s
 Wall time: 1h 12min 53s

In [43]:

```
%%time
model = AdaBoostClassifier(n_estimators = gridsrch.best_params_['n_estimators'])
model.fit(X_train, y_train)
plot_confusion_matrix(y_train, model.predict(X_train), 'Training')
print("The ROC-AUC Score obtained on CV set is : " ,roc_auc_score(y_cv, model.predict(X_cv)))
print("The F1 scores of each class on CV set are : " ,f1_score(y_cv, model.predict(X_cv),average='macro'))
print("The Macro F1-Score obtained on CV set is : " , f1_score(y_cv, model.predict(X_cv),average='macro'))
plot_confusion_matrix(y_cv, model.predict(X_cv), 'Cross Validation')
accuracy_check(model,X_test,y_test)
```


----- Training Confusion matrix -----




The ROC-AUC Score obtained on CV set is : 0.5350098131120029
 The F1 scores of each class on CV set are : [0.99646858 0.12080537]

The Macro F1-Score obtained on CV set is : 0.5586369751413783

----- Cross Validation Confusion matrix -----



----- Test Confusion matrix -----





The F1 scores : [0.99638023 0.08489097]

The Macro F1 score : 0.5406355964190894

CPU times: user 27min 46s, sys: 57.2 s, total: 28min 43s

Wall time: 28min 43s



Custom Ensemble on Robust scaled Data

In [44]:

```
class Custom_ensemble_model:

    def __init__(self, base_model, meta_model, number_of_samples, sample_size):

        self.base_model = base_model
        self.meta_model = meta_model
```

```

self.number_of_samples = number_of_samples
self.sample_size = sample_size

def base_model_split(self, D1_X, D1_y):

    X_sub_samples = []
    y_sub_samples = []

    #Run over all the subsamples
    for every_subsample in range(self.number_of_samples):
        #Create single subsample
        single_subsample = []
        possibilities = [0,1]
        #Create balanced dataset.
        for single_possibility in possibilities:
            samples = np.where(D1_y.values==single_possibility)[0]
            index_range = range(samples.shape[0])
            subsample_indexes = np.random.choice(index_range, size=self.sample_size, replace=True)
            single_subsample.extend(samples[subsample_indexes])
        X_sub_samples.append(D1_X[single_subsample])
        y_sub_samples.append(D1_y.values[single_subsample])

    return X_sub_samples, y_sub_samples

def train_base_models(self, X_sub_samples, y_sub_samples):

    base_models_train_list = []

    for i in range(len(X_sub_samples)):
        X_sub_samples[i], y_sub_samples[i] = shuffle(X_sub_samples[i], y_sub_samples[i], random_state=i)
        base_models_train_list.append(self.base_model.fit(X_sub_samples[i], y_sub_samples[i]))

    return base_models_train_list

def meta_model_dataset(self, base_models_train_list, D2_X, D2_y):

    meta_model_X = []
    meta_model_y = []

    for row in range(D2_X.shape[0]):
        for base_model in range(len(base_models_train_list)):
            prediction = base_models_train_list[base_model].predict(D2_X[row].reshape(1,-1))
            meta_model_X.append(prediction)
            meta_model_y.append(np.array(D2_y.values[row]))
    return meta_model_X, meta_model_y

def train_meta_model(self, meta_model_X, meta_model_y):

    final_model = self.meta_model.fit(meta_model_X, meta_model_y)
    print("ROC-AUC score of meta model : ", roc_auc_score(meta_model_y, final_model.predict(meta_model_X)))
    return final_model

```

In [45]:

```

%%time
base_model = DecisionTreeClassifier()
meta_model = XGBClassifier(n_estimators=300, nthread=-1, n_jobs=-1)

DT1_X, DT2_X, DT1_y, DT2_y = train_test_split(X_train, y_train, random_state = 42, stratify=y_train)

custom_ensemble_model = Custom_ensemble_model(base_model, meta_model, 100, 10000)
X_subsamples_list, y_subsamples_list = custom_ensemble_model.base_model_split(DT1_X, DT1_y)
base_models = custom_ensemble_model.train_base_models(X_subsamples_list, y_subsamples_list)
meta_model_X, meta_model_y = custom_ensemble_model.meta_model_dataset(base_models, DT2_X, DT2_y)
meta_model_final = custom_ensemble_model.train_meta_model(np.array(meta_model_X).reshape(-1,1),
meta_model_X, meta_model_y = [], [])

# test_X_remove, test_X, test_y_remove, test_y = train_test_split(X_test, y_test, random_state


```

```
# test_X_custom_model, test_y_custom_model = custom_ensemble_model.meta_model_dataset(base_model)
# accuracy_check(meta_model_final, np.array(test_X_custom_model).reshape(-1,1), np.array(test_y_c
```



[05:15:29] WARNING: ../src/learner.cc:1095: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.
ROC-AUC score of meta model : 0.8576382788512289
CPU times: user 23min 30s, sys: 12.2 s, total: 23min 42s
Wall time: 15min 16s

In [46]:

```
%%time
test_X_remove, test_X, test_y_remove, test_y = train_test_split(X_test, y_test, random_state =
test_X_custom_model, test_y_custom_model = custom_ensemble_model.meta_model_dataset(base_model)
accuracy_check(meta_model_final, np.array(test_X_custom_model).reshape(-1,1), np.array(test_y_c
```



The F1 scores : [0.99692308 0.]
The Macro F1 score : 0.49846153846153846
CPU times: user 51.8 s, sys: 39.2 ms, total: 51.8 s
Wall time: 51.4 s



Decision Tree Classifier on Log transformed and standard scaled Data

In [47]:

```
%%time
model = DecisionTreeClassifier()
parameters = {
    'criterion' : ['gini', 'entropy'],
    'max_depth' : [2,4,6,8,10,12]
}
clf = RandomizedSearchCV(model, parameters, scoring = 'roc_auc', n_jobs=-1)
```

```

gridsrch = clf.fit(X_train_log, y_train)
print("Best Params : " , gridsrch.best_params_)
print("Best Score : " , gridsrch.best_score_)

```

```

Best Params : {'max_depth': 8, 'criterion': 'entropy'}
Best Score : 0.9316775985007787
CPU times: user 7.67 s, sys: 761 ms, total: 8.43 s
Wall time: 1min 7s

```

In [48]:


```

%%time
model = DecisionTreeClassifier(
    criterion = gridsrch.best_params_['criterion'],
    max_depth = gridsrch.best_params_['max_depth'])
model.fit(X_train_log, y_train)
plot_confusion_matrix(y_train, model.predict(X_train_log), 'Training')
print("The ROC-AUC Score obtained on CV set is : " , roc_auc_score(y_cv, model.predict(X_cv_log)))
print("The F1 scores of each class on CV set are : " , f1_score(y_cv, model.predict(X_cv_log), average='macro'))
print("The Macro F1-Score obtained on CV set is : " , f1_score(y_cv, model.predict(X_cv_log), average='macro'))
plot_confusion_matrix(y_cv, model.predict(X_cv_log), 'Cross Validation')

accuracy_check(model,X_test_log,y_test)

```

----- Training Confusion matrix -----




```


The ROC-AUC Score obtained on CV set is : 0.5010875760267487
The F1 scores of each class on CV set are : [0.99654363 0.00443459]
The Macro F1-Score obtained on CV set is : 0.5004891089181579

```



----- Cross Validation Confusion matrix -----



----- Test Confusion matrix -----



The F1 scores : [0.99655744 0.01236749]
 The Macro F1 score : 0.5044624638902091
 CPU times: user 21.8 s, sys: 28 ms, total: 21.9 s
 Wall time: 21.8 s



Random Forest Classifier on Log transformed and standard scaled Data

In [49]:

```
%%time
model = RandomForestClassifier(class_weight = "balanced_subsample" , n_jobs = -1)
parameters = {
    'n_estimators' : [10,50,100,300,500,1000],
    'max_depth' : [1,3,5,7]
}
clf = RandomizedSearchCV(model, parameters, scoring = 'roc_auc', n_jobs=-1)
gridsrch = clf.fit(X_train_log, y_train)
print("Best Params : " , gridsrch.best_params_)
print("Best Score : " , gridsrch.best_score_)
```

Best Params : {'n_estimators': 1000, 'max_depth': 7}
 Best Score : 0.9368956651927901
 CPU times: user 41min 11s, sys: 11.3 s, total: 41min 23s
 Wall time: 1h 21min 1s

In [50]:

```
%%time
model = RandomForestClassifier(
    n_estimators = gridsrch.best_params_['n_estimators'],
    max_depth = gridsrch.best_params_['max_depth'],
    class_weight = "balanced_subsample",
    n_jobs = -1)
model.fit(X_train_log, y_train)
plot_confusion_matrix(y_train, model.predict(X_train_log), 'Training')
```

```

print("The ROC-AUC Score obtained on CV set is : ",roc_auc_score(y_cv, model.predict(X_cv_log),ave
print("The F1 scores of each class on CV set are : ",f1_score(y_cv, model.predict(X_cv_log),ave
print("The Macro F1-Score obtained on CV set is : ", f1_score(y_cv, model.predict(X_cv_log),av
plot_confusion_matrix(y_cv, model.predict(X_cv_log), 'Cross Validation')


```

```

accuracy_check(model,X_test_log,y_test)

```

----- Training Confusion matrix -----




The ROC-AUC Score obtained on CV set is : 0.8686192243369859


The F1 scores of each class on CV set are : [0.92277379 0.07783183]

The Macro F1-Score obtained on CV set is : 0.5003028097328748

----- Cross Validation Confusion matrix -----



----- Test Confusion matrix -----





The F1 scores : [0.92155153 0.07544431]

The Macro F1 score : 0.4984979177183212

CPU times: user 46min 48s, sys: 12.2 s, total: 47min 1s

Wall time: 8min 3s



Balanced Bagging classifier on Log transformed and standard scaled Data

In [51]:

```
model = BalancedBaggingClassifier(n_jobs = -1)
parameters = {'n_estimators' : [10,50,100,300,500,1000]}
clf = RandomizedSearchCV(model, parameters,scoring = 'roc_auc', n_jobs=-1)
gridsrch = clf.fit(X_train_log, y_train)
print("Best Params : " , gridsrch.best_params_)
print("Best Score : " , gridsrch.best_score_)
```


Best Params : {'n_estimators': 1000}
Best Score : 0.9684490921198463

In [52]:

```
model = BalancedBaggingClassifier(
    n_estimators = gridsrch.best_params_['n_estimators'],
    n_jobs = -1)
model.fit(X_train_log, y_train)
plot_confusion_matrix(y_train, model.predict(X_train_log), 'Training')
print("The ROC-AUC Score obtained on CV set is : " ,roc_auc_score(y_cv, model.predict(X_cv_log))
print("The F1 scores of each class on CV set are : " ,f1_score(y_cv, model.predict(X_cv_log),ave
print("The Macro F1-Score obtained on CV set is : " , f1_score(y_cv, model.predict(X_cv_log),ave
plot_confusion_matrix(y_cv, model.predict(X_cv_log), 'Cross Validation')

accuracy_check(model,X_test_log,y_test)
```

----- Training Confusion matrix -----




The ROC-AUC Score obtained on CV set is : 0.9195372798779126


The F1 scores of each class on CV set are : [0.9567805 0.13268687]

The Macro F1-Score obtained on CV set is : 0.5447336862994457

----- Cross Validation Confusion matrix -----





----- Test Confusion matrix -----



The F1 scores : [0.95636144 0.13148611]

The Macro F1 score : 0.5439237793562582



XGBoost classifier on Log transformed and standard scaled Data

In [41]:

```
# %%time
# model = XGBClassifier(nthread=-1, n_jobs=-1)
# parameters = {'n_estimators' : [10,50,100,300,500,1000]}
# clf = RandomizedSearchCV(model, parameters,scoring = 'roc_auc', n_jobs=-1)
# gridsrch = clf.fit(X_train_log, y_train)
# print("Best Params : " , gridsrch.best_params_)
# print("Best Score : " , gridsrch.best_score_)
```

```
[07:41:05] WARNING: ../src/learner.cc:1095: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.
Best Params : {'n_estimators': 1000}
Best Score : 0.972790999354252
CPU times: user 54min 59s, sys: 1.09 s, total: 55min
Wall time: 2h 2min 16s
```

In [53]:

```
import random
from tqdm import tqdm
def custom_loops(x_train,y_train,classifier, param_range, folds):
    testscores = []
    #Referance - https://www.geeksforgeeks.org/python-random-sample-function/
    #1.generate 10 unique values(uniform random distribution) in the given range "param_range"

    lst = random.sample(param_range,5)
    lst.sort()

    params = {'n_estimators':lst}

    for k in tqdm(params['n_estimators']):
        testscores_folds = []
        for fold in range(0, folds):
            #2.devide numbers ranging from 0 to len(X_train) into groups= folds
            block_size = int(len(x_train)/folds)
            test_indices = list(set(list(range((block_size*fold), (block_size*(fold+1))))))
            train_indices = list(set(list(range(1, len(x_train)))) - set(test_indices))
            X_train = pd.DataFrame(x_train).iloc[train_indices]
            Y_train = pd.DataFrame(y_train).iloc[train_indices]
            X_test = pd.DataFrame(x_train).iloc[test_indices]
            Y_test = pd.DataFrame(y_train).iloc[test_indices]

            classifier.n_estimators = k
            classifier.fit(X_train,Y_train)

            Y_predicted = classifier.predict(X_test)
            testscores_folds.append(f1_score(Y_test, Y_predicted, average = 'macro'))

        testscores.append(np.mean(np.array(testscores_folds)))
    return testscores,params
```

In [54]:

```
%%time
model = XGBClassifier(nthread=-1)
parameters = [10,50,100,300,500,1000]
#clf = RandomizedSearchCV(model, parameters,scoring = 'roc_auc', n_jobs=-1)
#gridsrch = clf.fit(X_train, y_train)

testscores,paramsopt = custom_loops(X_train_log, y_train, model, parameters, 3)
print("Params : " , paramsopt)
print("Scores : " , testscores)
```

```
0%|          | 0/5 [00:00<?, ?it/s]
[07:33:21] WARNING: ../src/learner.cc:1095: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.
[07:33:35] WARNING: ../src/learner.cc:1095: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.
[07:33:50] WARNING: ../src/learner.cc:1095: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.
20%|██        | 1/5 [00:44<02:58, 44.70s/it]
[07:34:05] WARNING: ../src/learner.cc:1095: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.
[07:34:32] WARNING: ../src/learner.cc:1095: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.
```

```

[07:34:59] WARNING: ../src/learner.cc:1095: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.
 40%|███████| 2/5 [02:05<03:17, 65.90s/it]
[07:35:25] WARNING: ../src/learner.cc:1095: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.
[07:36:45] WARNING: ../src/learner.cc:1095: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.
[07:38:04] WARNING: ../src/learner.cc:1095: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.
 60%|███████| 3/5 [06:04<04:50, 145.07s/it]
[07:39:25] WARNING: ../src/learner.cc:1095: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.
[07:41:34] WARNING: ../src/learner.cc:1095: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.
[07:43:32] WARNING: ../src/learner.cc:1095: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.
 80%|███████| 4/5 [12:08<03:51, 231.30s/it]
[07:45:28] WARNING: ../src/learner.cc:1095: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.
[07:49:16] WARNING: ../src/learner.cc:1095: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.
[07:53:13] WARNING: ../src/learner.cc:1095: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.
100%|██████████| 5/5 [23:41<00:00, 284.31s/it]
Params : {'n_estimators': [50, 100, 300, 500, 1000]}
Scores : [0.5462730824495994, 0.5709103739719259, 0.6317776680594849, 0.6606885765951587, 0.699680425280849]
CPU times: user 3h 7min 49s, sys: 2.4 s, total: 3h 7min 52s
Wall time: 23min 41s

```

In [55]:

```

%%time
model = XGBClassifier(
    n_estimators = 1000,
    nthread = -1,
    n_jobs=-1)
model.fit(X_train_log, y_train)
plot_confusion_matrix(y_train, model.predict(X_train_log), 'Training')
print("The ROC-AUC Score obtained on CV set is : ", roc_auc_score(y_cv, model.predict(X_cv_log)))
print("The F1 scores of each class on CV set are : ", f1_score(y_cv, model.predict(X_cv_log), average='micro'))
print("The Macro F1-Score obtained on CV set is : ", f1_score(y_cv, model.predict(X_cv_log), average='macro'))
plot_confusion_matrix(y_cv, model.predict(X_cv_log), 'Cross Validation')


accuracy_check(model, X_test_log, y_test)

```

```

[07:57:46] WARNING: ../src/learner.cc:1095: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.
----- Training Confusion matrix -----
-----

```




The ROC-AUC Score obtained on CV set is : 0.6574755784001527


The F1 scores of each class on CV set are : [0.99743723 0.45802771]

The Macro F1-Score obtained on CV set is : 0.7277324706503238

----- Cross Validation Confusion matrix -----



----- Test Confusion matrix -----





The F1 scores : [0.99646902 0.03857442]

The Macro F1 score : 0.5175217223229691

CPU times: user 48min 11s, sys: 470 ms, total: 48min 12s

Wall time: 6min 22s



Adaboost classifier on Log transformed and standard scaled Data

In [56]:

```
%%time
model = AdaBoostClassifier()
parameters = {'n_estimators' : [10,50,100,300,500,1000]}
clf = RandomizedSearchCV(model, parameters,scoring = 'roc_auc', n_jobs=-1)
gridsrch = clf.fit(X_train_log, y_train)
print("Best Params : " , gridsrch.best_params_)
print("Best Score : " , gridsrch.best_score_)
```


Best Params : {'n_estimators': 1000}
 Best Score : 0.9399416563817754
 CPU times: user 22min 55s, sys: 952 ms, total: 22min 56s
 Wall time: 1h 8min 26s

In [57]:

```
%%time
model = AdaBoostClassifier(n_estimators = gridsrch.best_params_['n_estimators'])
model.fit(X_train_log, y_train)
plot_confusion_matrix(y_train, model.predict(X_train_log), 'Training')
print("The ROC-AUC Score obtained on CV set is : " ,roc_auc_score(y_cv, model.predict(X_cv_log))
print("The F1 scores of each class on CV set are : " ,f1_score(y_cv, model.predict(X_cv_log),average='macro')
print("The Macro F1-Score obtained on CV set is : " , f1_score(y_cv, model.predict(X_cv_log),average='macro')
plot_confusion_matrix(y_cv, model.predict(X_cv_log), 'Cross Validation')

accuracy_check(model,X_test_log,y_test)
```

----- Training Confusion matrix -----




The ROC-AUC Score obtained on CV set is : 0.534425473598223


The F1 scores of each class on CV set are : [0.99644156 0.11832061]

The Macro F1-Score obtained on CV set is : 0.5573810829638208

----- Cross Validation Confusion matrix -----



----- Test Confusion matrix -----





The F1 scores : [0.99612914 0.11173975]

The Macro F1 score : 0.5539344445034765

CPU times: user 29min 48s, sys: 442 ms, total: 29min 49s

Wall time: 29min 49s



MLP on Robust scaled Data

In [64]:

```
import keras
from keras.models import Sequential
from keras.layers import Dense
from keras import initializers
import tensorflow as tf
```

```
from sklearn.metrics import roc_auc_score
```

```
metrics = [  
    keras.metrics.AUC(name='roc_auc'),  
    keras.metrics.AUC(name='pr_auc', curve='PR'),  
]
```

In [65]:

```
mlp_model = Sequential()  
mlp_model.add(Dense(15, input_dim=21, activation='tanh'))  
mlp_model.add(Dense(1, activation='sigmoid'))  
mlp_model.compile(loss='binary_crossentropy', optimizer='adam', metrics=metrics)  
  
mlp_model_history = mlp_model.fit(X_train, y_train, epochs=10, batch_size=64, validation_data=(
```

```
Epoch 1/10  
18332/18332 [=====] - 23s 1ms/step - loss: 0.0500 - roc_auc: 0.8130 -  
pr_auc: 0.0441 - val_loss: 0.0264 - val_roc_auc: 0.9218 - val_pr_auc: 0.1852  
Epoch 2/10  
18332/18332 [=====] - 23s 1ms/step - loss: 0.0268 - roc_auc: 0.9157 -  
pr_auc: 0.1795 - val_loss: 0.0262 - val_roc_auc: 0.9302 - val_pr_auc: 0.1934  
Epoch 3/10  
18332/18332 [=====] - 23s 1ms/step - loss: 0.0262 - roc_auc: 0.9191 -  
pr_auc: 0.1822 - val_loss: 0.0260 - val_roc_auc: 0.9280 - val_pr_auc: 0.1916  
Epoch 4/10  
18332/18332 [=====] - 23s 1ms/step - loss: 0.0263 - roc_auc: 0.9204 -  
pr_auc: 0.1843 - val_loss: 0.0260 - val_roc_auc: 0.9252 - val_pr_auc: 0.1876  
Epoch 5/10  
18332/18332 [=====] - 23s 1ms/step - loss: 0.0263 - roc_auc: 0.9222 -  
pr_auc: 0.1912 - val_loss: 0.0261 - val_roc_auc: 0.9270 - val_pr_auc: 0.1900  
Epoch 6/10  
18332/18332 [=====] - 23s 1ms/step - loss: 0.0265 - roc_auc: 0.9204 -  
pr_auc: 0.1907 - val_loss: 0.0259 - val_roc_auc: 0.9283 - val_pr_auc: 0.1930  
Epoch 7/10  
18332/18332 [=====] - 23s 1ms/step - loss: 0.0259 - roc_auc: 0.9256 -  
pr_auc: 0.1881 - val_loss: 0.0259 - val_roc_auc: 0.9233 - val_pr_auc: 0.1971  
Epoch 8/10  
18332/18332 [=====] - 23s 1ms/step - loss: 0.0265 - roc_auc: 0.9203 -  
pr_auc: 0.1913 - val_loss: 0.0259 - val_roc_auc: 0.9277 - val_pr_auc: 0.1975  
Epoch 9/10  
18332/18332 [=====] - 23s 1ms/step - loss: 0.0259 - roc_auc: 0.9182 -  
pr_auc: 0.1923 - val_loss: 0.0259 - val_roc_auc: 0.9305 - val_pr_auc: 0.1970  
Epoch 10/10  
18332/18332 [=====] - 23s 1ms/step - loss: 0.0260 - roc_auc: 0.9227 -  
pr_auc: 0.1975 - val_loss: 0.0259 - val_roc_auc: 0.9300 - val_pr_auc: 0.1964
```

In [66]:

```
print(mlp_model.summary())
```

Model: "sequential_5"


Layer (type)	Output Shape	Param #
dense_10 (Dense)	(None, 15)	330
dense_11 (Dense)	(None, 1)	16

Total params: 346
Trainable params: 346
Non-trainable params: 0

None

In [67]:


```
plot.plot(mlp_model_history.history['roc_auc'])  
plot.plot(mlp_model_history.history['val_roc_auc'])  
plot.title('Accuracy Plot')  
plot.ylabel('Accuracy')  
plot.xlabel('Epoch')  
plot.legend(['Train', 'Validation'])  
plot.show()
```



```
In [69]: mlp_results = mlp_model.evaluate(X_test, y_test, batch_size=64, verbose=0)
for name, value in zip(mlp_model.metrics_names, mlp_results):
    print(name, ': ', value)
```

```
loss : 0.026836572214961052
roc_auc : 0.9172110557556152
pr_auc : 0.17135468125343323
```

```
In [70]: mlp_test_pred_15 = mlp_model.predict_classes(X_test)
plot_confusion_matrix(y_test, mlp_test_pred_15, 'Test')
```



```
In [71]: # Reference - http://zetcode.com/python/prettytable/
from prettytable import PrettyTable

x = PrettyTable()

x.field_names = ["Data", "Model", "Test Precision Recall AUC", "Test ROC-AUC"]

x.add_row(["Robust Scaled Data", "Decision Tree", 0.0174, 0.748])
x.add_row(["Robust Scaled Data", "Random Forest", 0.161, 0.932])
x.add_row(["Robust Scaled Data", "Balanced Bagging Classifier", 0.205, 0.708])
x.add_row(["Robust Scaled Data", "XGBoost Classifier", 0.553, 0.861])
x.add_row(["Robust Scaled Data", "AdaBoost Classifier", 0.164, 0.524])
x.add_row(["Robust Scaled Data", "Custom Ensemble Model", 0.003, 0.5])
x.add_row(["Robust Scaled Data", "Perceptron Model", 0.171, 0.917])
x.add_row(["Log Transformed Standard Scaled Data", "Decision Tree", 0.157, 0.625])
x.add_row(["Log Transformed Standard Scaled Data", "Random Forest", 0.176, 0.932])
x.add_row(["Log Transformed Standard Scaled Data", "Balanced Bagging Classifier", 0.208, 0.71])
x.add_row(["Log Transformed Standard Scaled Data", "XGBoost Classifier", 0.086, 0.57])
x.add_row(["Log Transformed Standard Scaled Data", "AdaBoost Classifier", 0.135, 0.535])

print(x)
```


C		Data		Model	Test Precision Recall AU
Test ROC-AUC					
0.748	Robust Scaled Data		Decision Tree		0.0174
0.932	Robust Scaled Data		Random Forest		0.161
0.708	Robust Scaled Data		Balanced Bagging Classifier		0.205
0.861	Robust Scaled Data		XGBoost Classifier		0.553
0.524	Robust Scaled Data		AdaBoost Classifier		0.164
0.5	Robust Scaled Data		Custom Ensemble Model		0.003
0.917	Robust Scaled Data		Perceptron Model		0.171
0.625	Log Transformed Standard Scaled Data		Decision Tree		0.157
0.932	Log Transformed Standard Scaled Data		Random Forest		0.176
0.71	Log Transformed Standard Scaled Data		Balanced Bagging Classifier		0.208
0.57	Log Transformed Standard Scaled Data		XGBoost Classifier		0.086
0.535	Log Transformed Standard Scaled Data		AdaBoost Classifier		0.135

In conclusion we can say for both **Robust Scaled and Log Transformed Data Random Forest** give us the best result which have **ROC-AUC Score 0.932**.

In []: